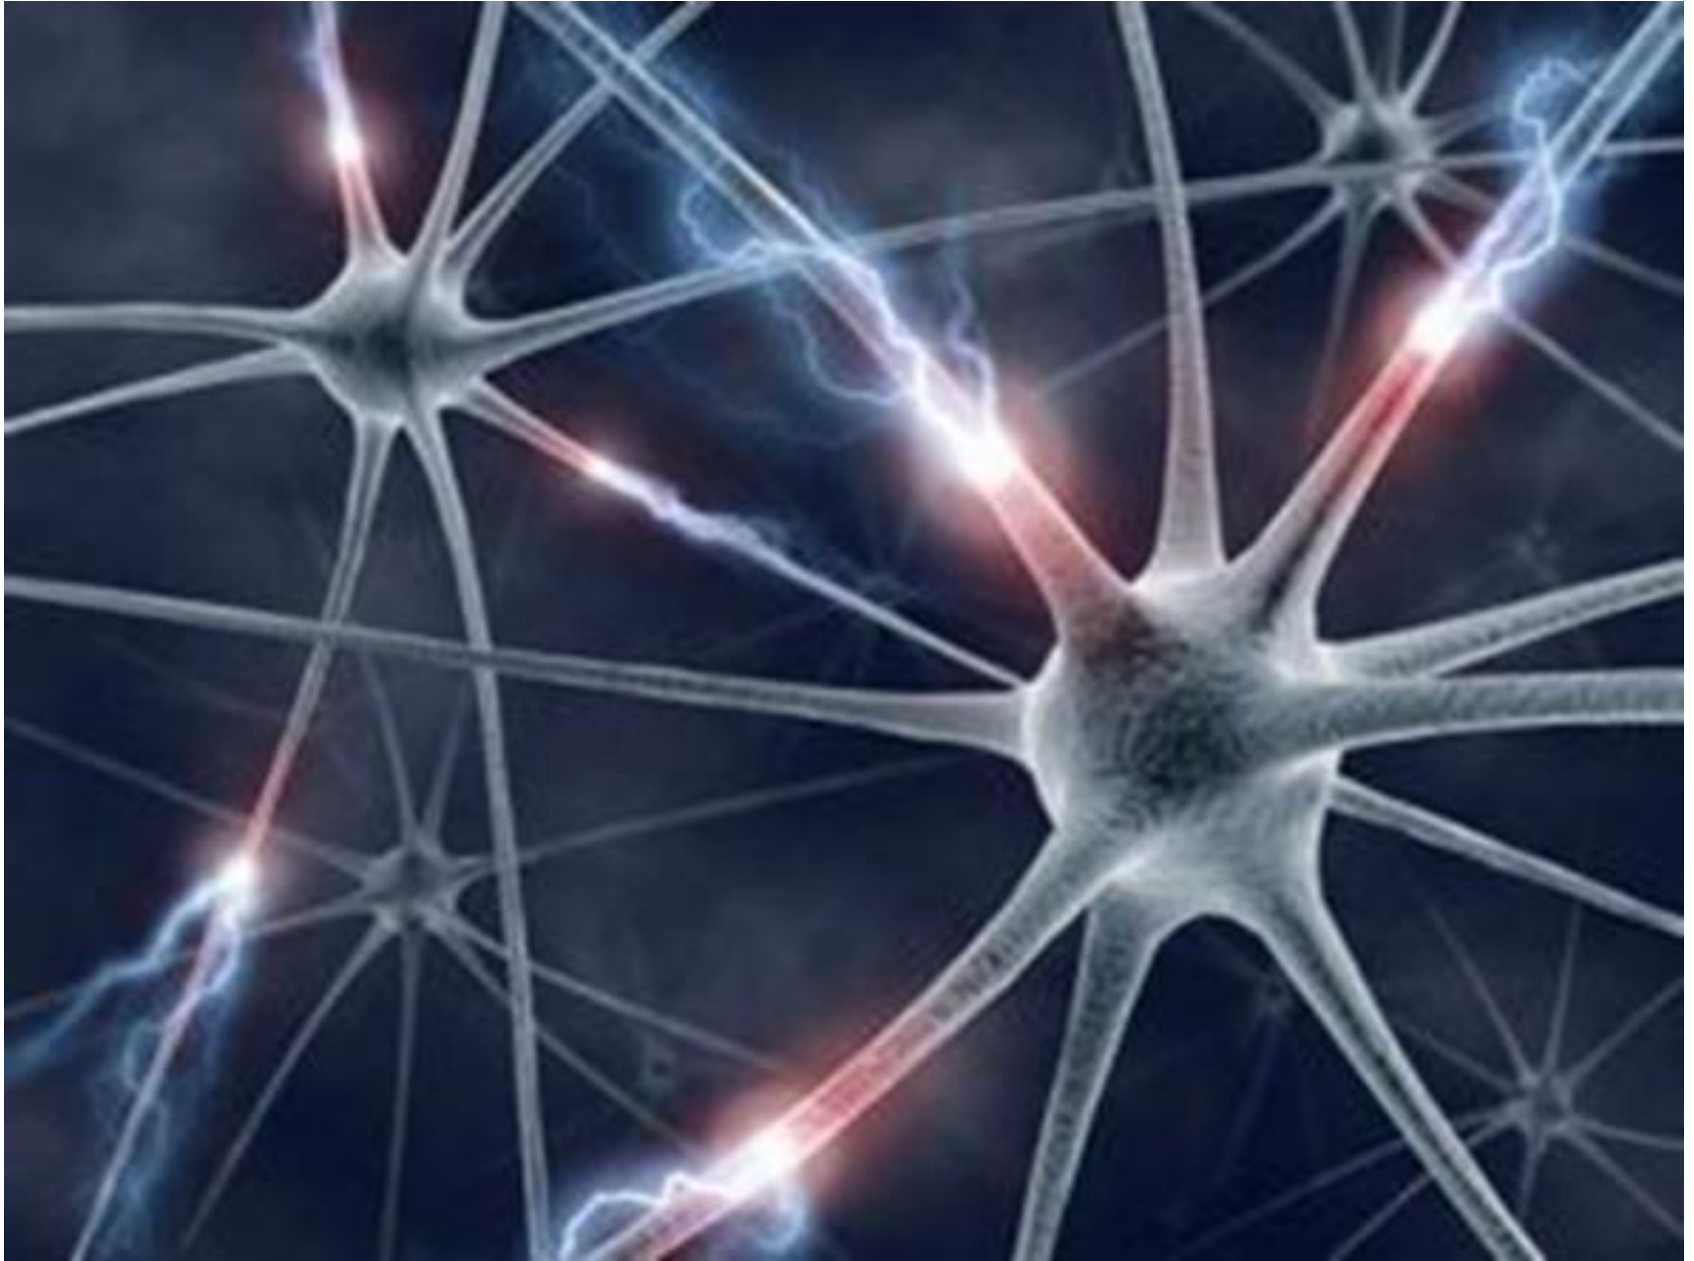


Neural Network

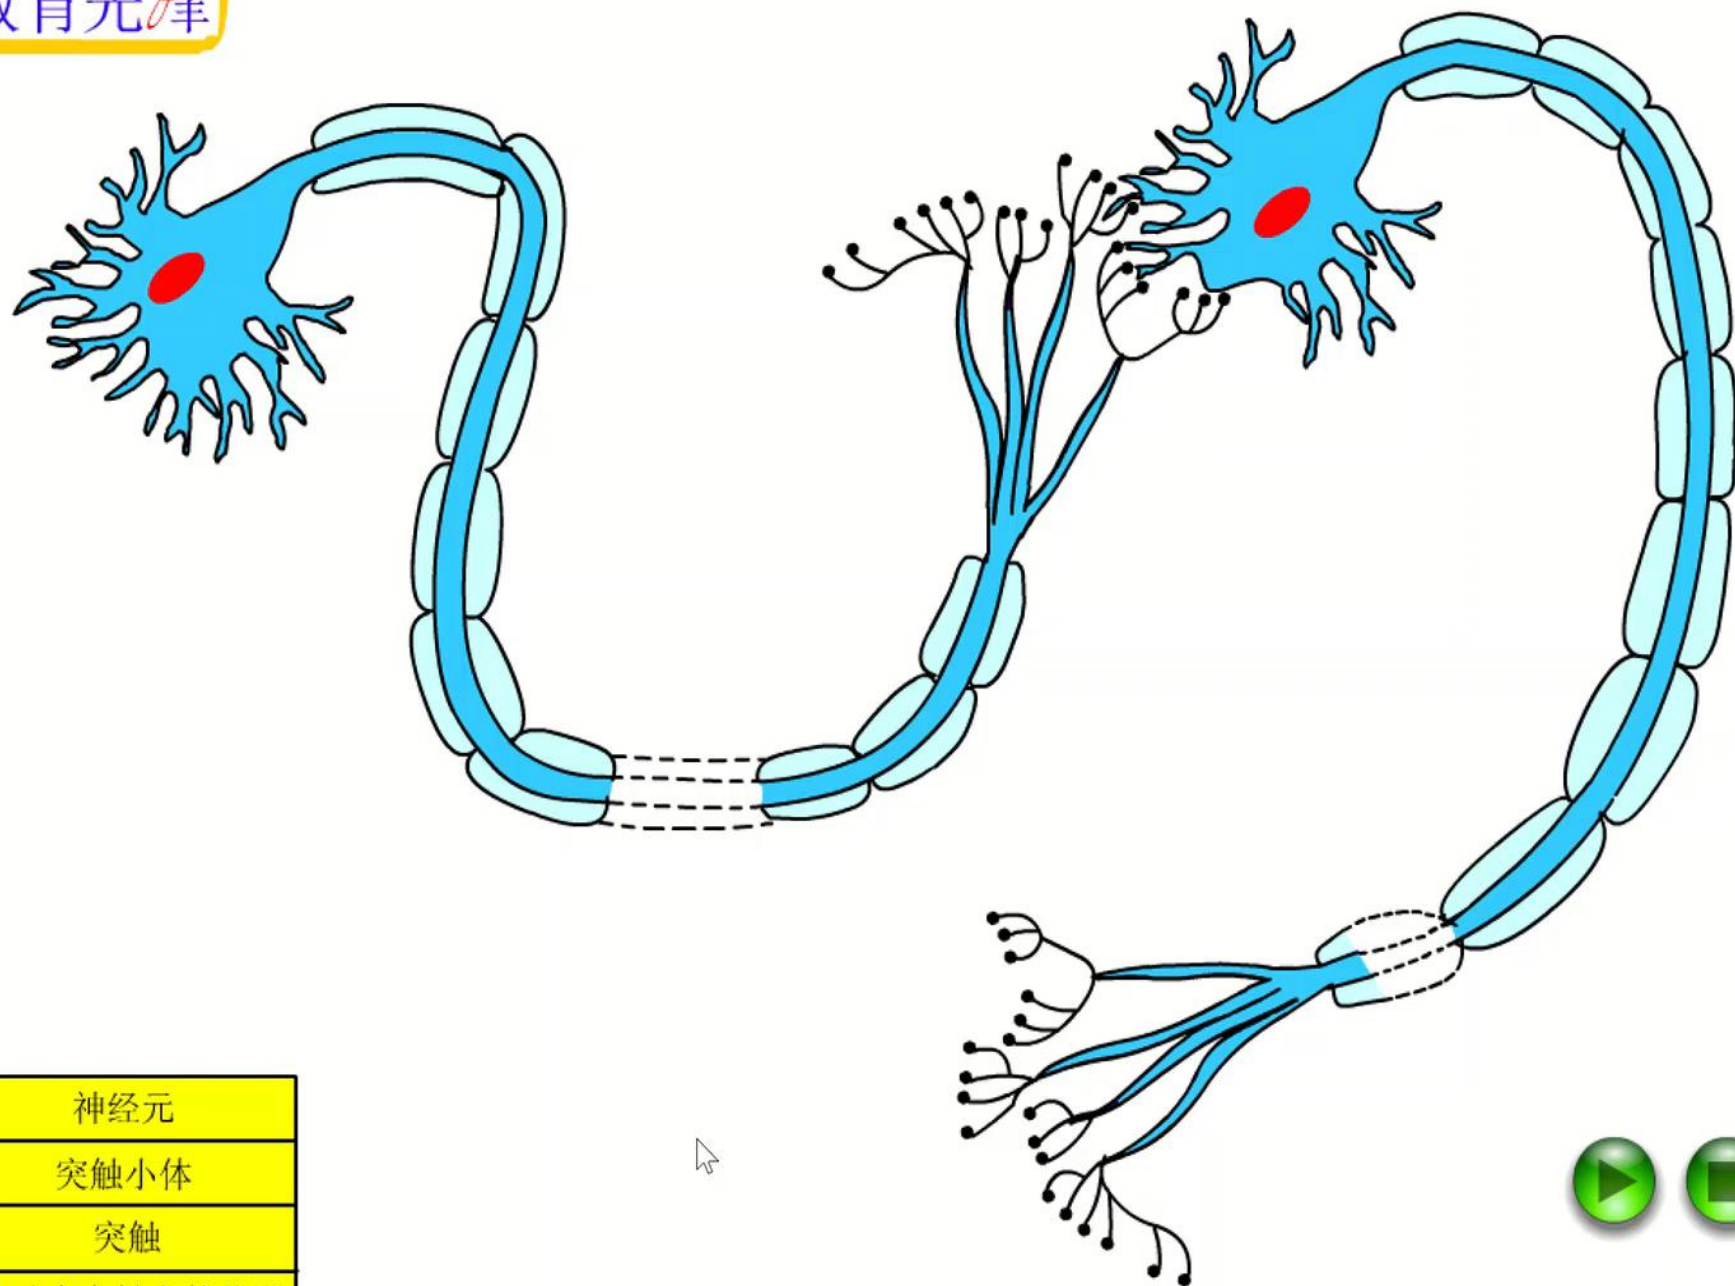


Outline

- ❖ Modeling one neuron
 - ❖ Activation functions
 - ❖ Fully connected feed-forward network
 - ❖ How to train a multi-layer network
 - ❖ Representational power of NN
-

Biology

- Neurons respond slowly
 - 10^{-3} s compared to 10^{-9} s for electrical circuits
 - The brain uses massively parallel computation
 - $\approx 10^{11}$ neurons in the brain
 - $\approx 10^4$ connections per neuron
-



神经元

突触小体

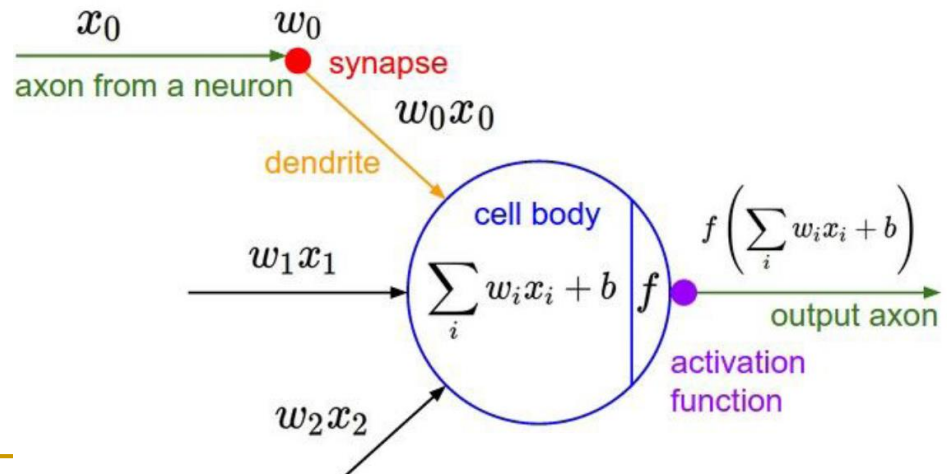
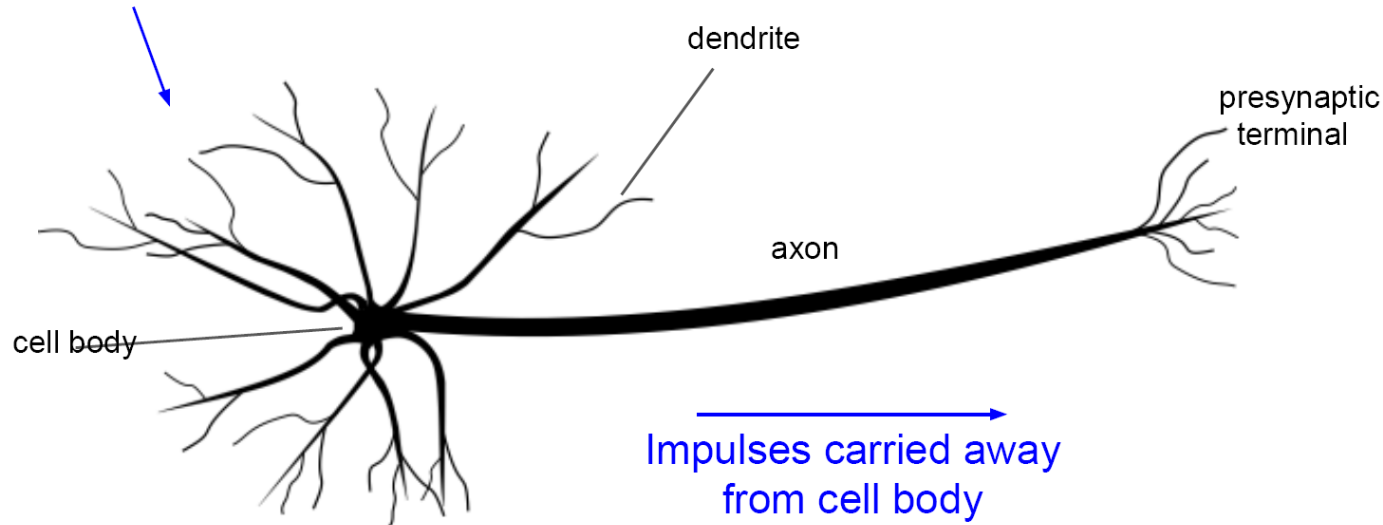
突触

冲动在突触上的传递

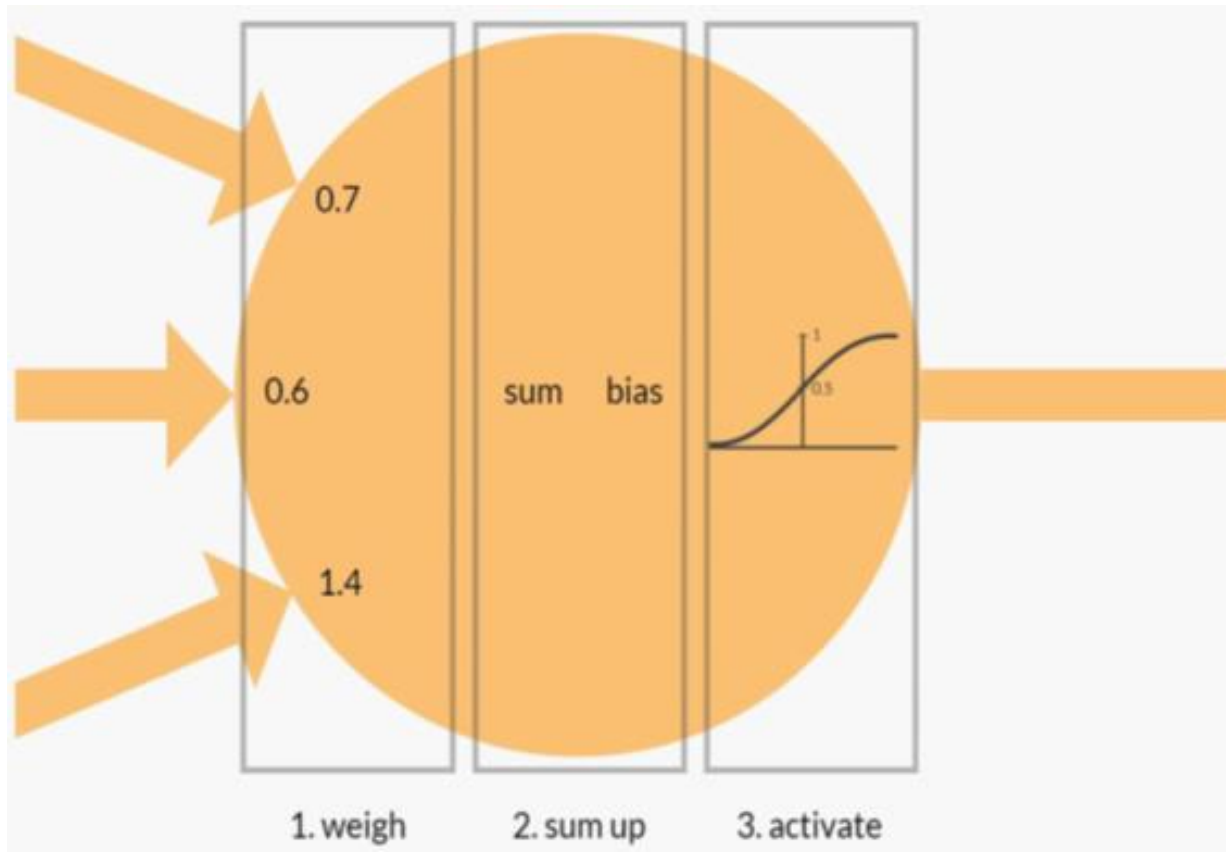


Modeling one neuron

Impulses carried toward cell body



Neuron



Single neuron as a linear classifier

- ❖ With an appropriate loss function on the neuron's output, a single neuron can be turned into a linear classifier.
- ❖ A single neuron can be used to implement a binary classifier (e.g. binary Softmax or binary SVM classifiers).

Single neuron as a linear classifier

❖ Binary Softmax classifier

◆ Interpreting:

- $\sigma(\sum_i w_i x_i + b)$ to be the probability of one of the classes

$$P(y_i = 1 \mid x_i; w)$$

- The probability of the other class to be

$$P(y_i = 0 \mid x_i; w) = 1 - P(y_i = 1 \mid x_i; w)$$

- ◆ With this interpretation, we can formulate the **cross-entropy loss**, and optimizing it would lead to a **binary Softmax classifier** (also known as logistic regression)

- ❖ Alternatively, we could attach a **max-margin hinge loss** to the output of the neuron and train it to become a **binary Support Vector Machine**.

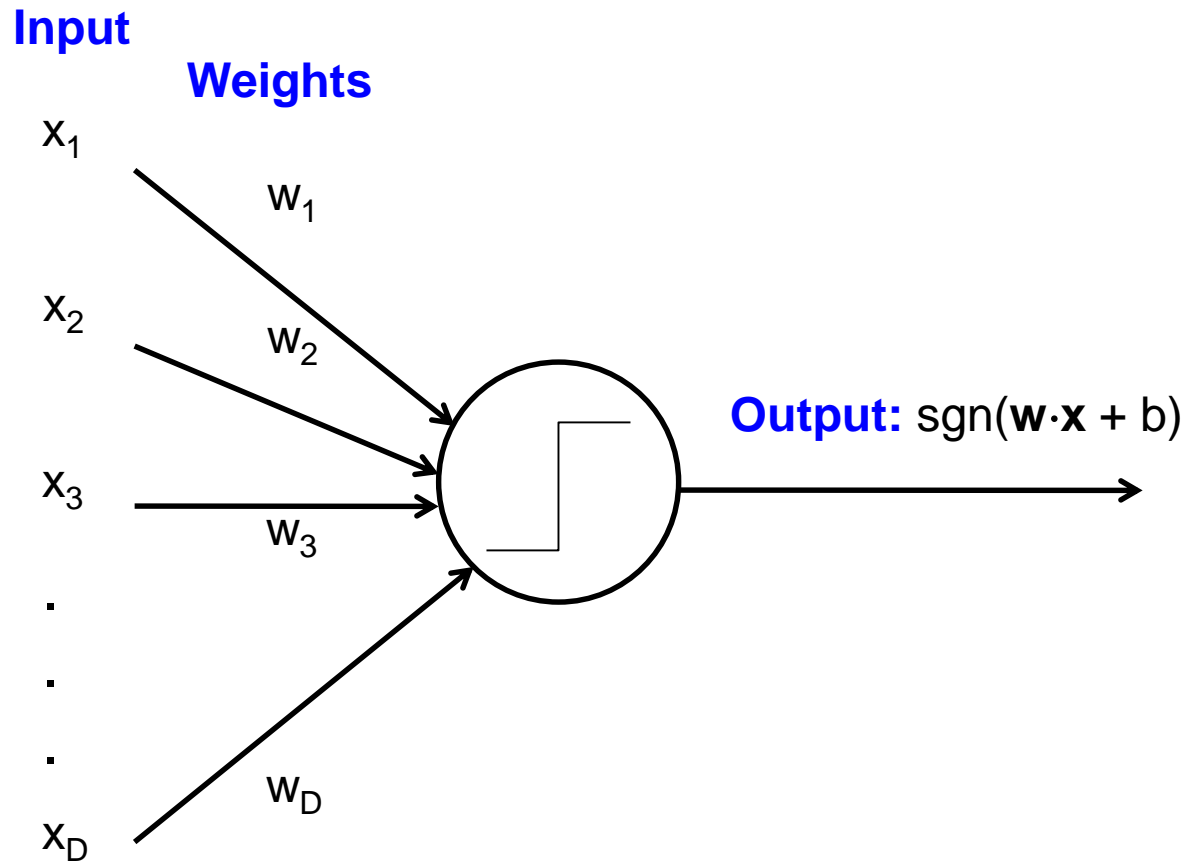
Outline

- ❖ Modeling one neuron
 - ❖ Activation functions
 - ❖ Fully connected feed-forward network
 - ❖ How to train a multi-layer network
 - ❖ Representational power of NN
-

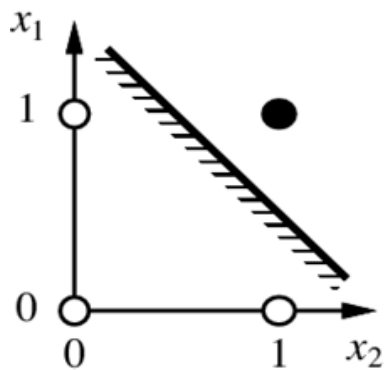
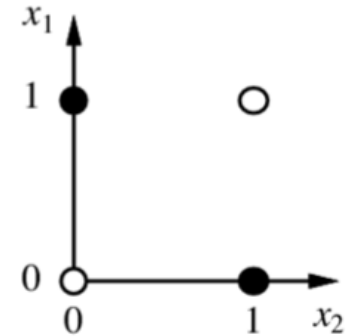
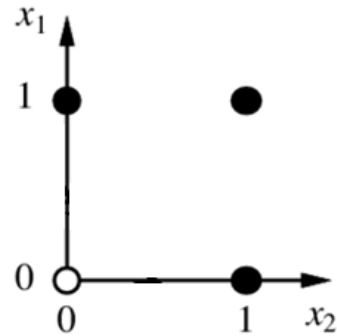
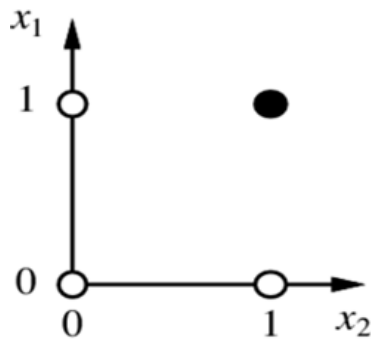
Activation Functions

- ❖ Neural networks are used to implement complex functions, and **non-linear** Activation Functions enable them to approximate arbitrarily complex functions.
- ❖ Without the non-linearity introduced by the Activation Function:
 - ◆ a neuron is equivalent to a linear classifier
 - ◆ a multi-layer neural network is equivalent to a single layer neural network

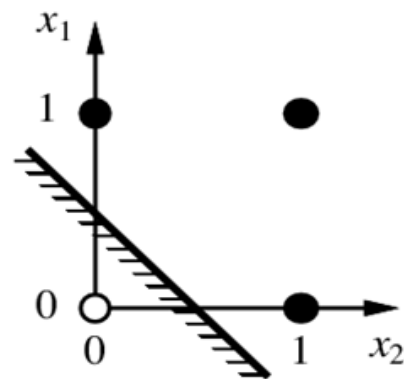
Perceptron



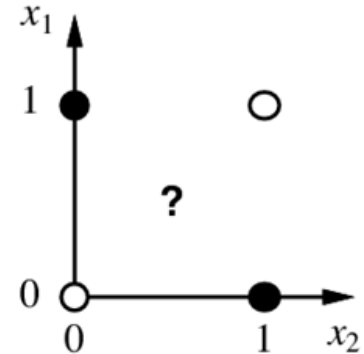
Linear separability



x_1 **and** x_2



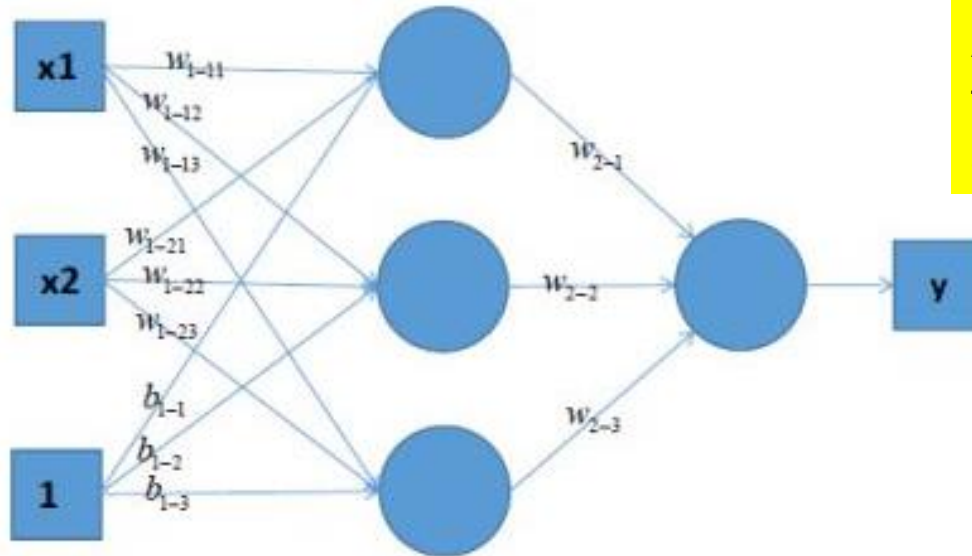
x_1 **or** x_2



x_1 **xor** x_2

Multilayer Perceptron (MLP)

❖ A MLP with a single hidden layer



Add nonlinearity
though
activation functions!

$$y = w_{2-1}(w_{1-11}x_1 + w_{1-21}x_2 + b_{1-1}) + w_{2-2}(w_{1-12}x_1 + w_{1-22}x_2 + b_{1-2}) + w_{2-3}(w_{1-13}x_1 + w_{1-23}x_2 + b_{1-3})$$



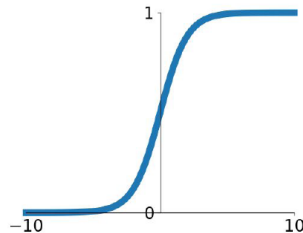
$$y = x_1(w_{2-1}w_{1-11} + w_{2-2}w_{1-12} + w_{2-3}w_{1-13}) + x_2(w_{2-1}w_{1-21} + w_{2-2}w_{1-22} + w_{2-3}w_{1-23}) + w_{2-1}b_{1-1} + w_{2-2}b_{1-2} + w_{2-3}b_{1-3}$$

Still a linear classifier

Some Activation Functions

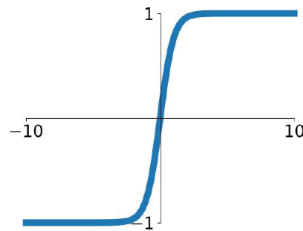
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



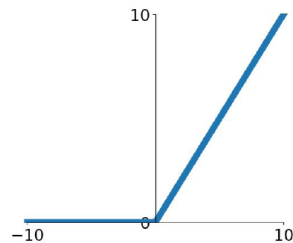
tanh

$$\tanh(x)$$



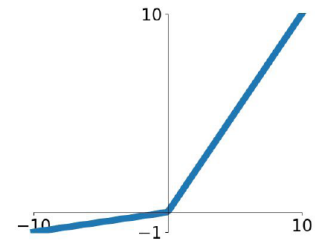
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

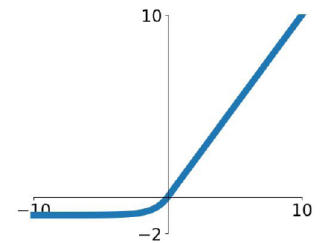


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



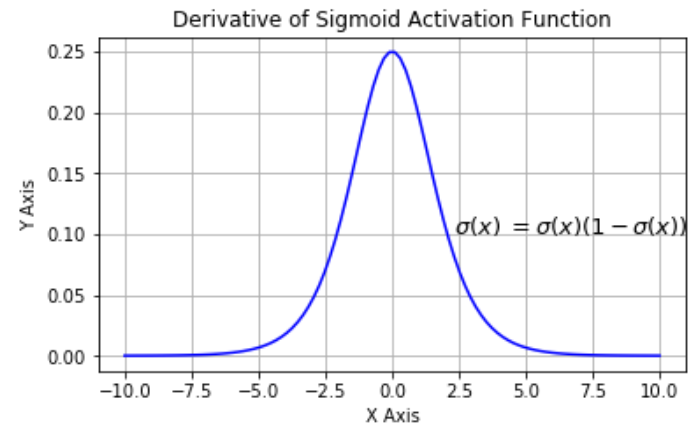
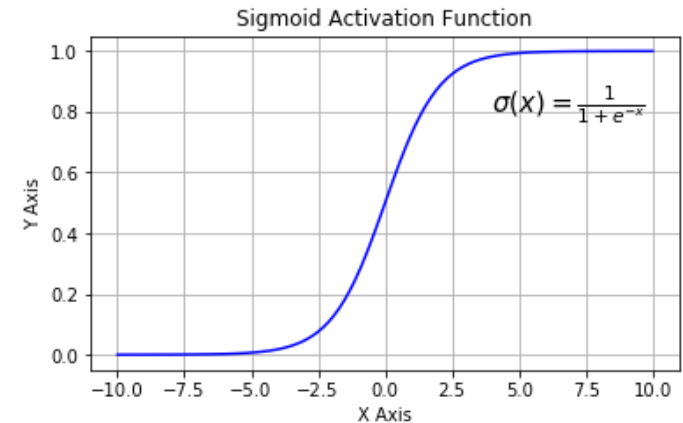
Activation Functions

❖ Sigmoid $\sigma(x) = 1/(1 + e^{-x})$

- ◆ Squashes numbers to range [0,1]
- ◆ Historically popular since they have nice interpretation as a saturating “firing rate” of a neuron

◆ Problems:

- ❑ Saturated neurons “kill” the gradients
- ❑ Sigmoid outputs are not zero-centered
- ❑ $\exp()$ is a bit compute expensive



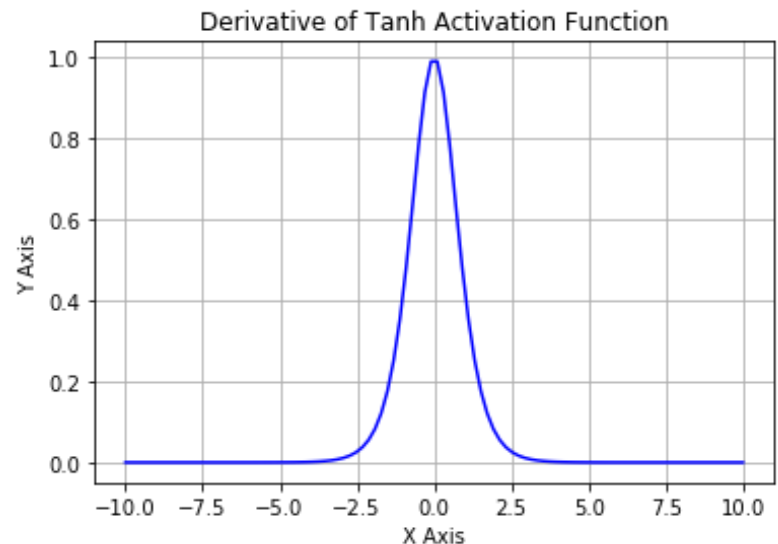
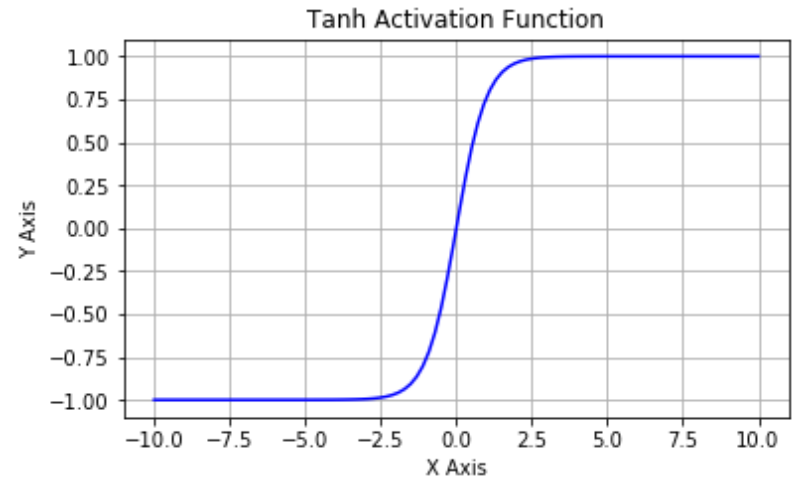
Activation Functions

$$\tanh x = \frac{\sinh x}{\cosh x} = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

❖ Tanh:

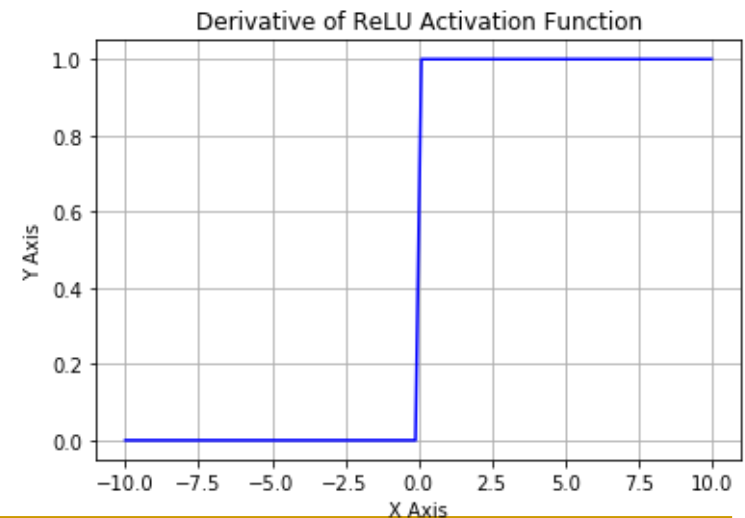
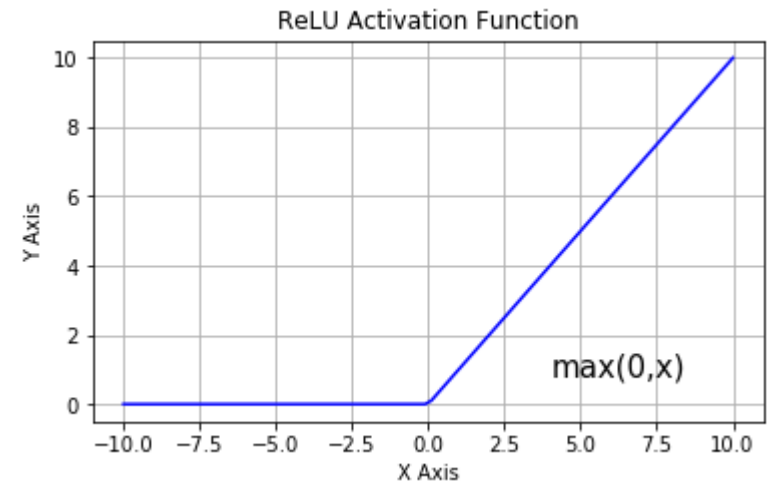
$$\tanh(x) = 2\sigma(2x) - 1$$

- ◆ Squashes numbers to range $[-1, 1]$
- ◆ Zero centered
- ◆ Problem:
 - ▣ kills gradients when saturated



Activation Functions

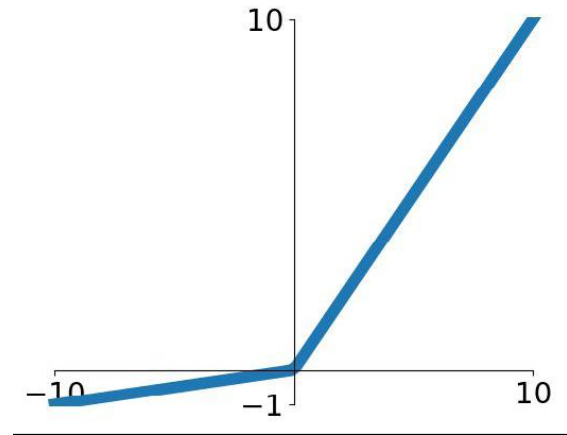
- ❖ **Relu: $f(x) = \max(0, x)$**
(Rectified Linear Unit)
 - ◆ Does not saturate (in +region)
 - ◆ Very computationally efficient
 - ◆ Converges much faster than sigmoid/tanh in practice (e.g. 6x)
 - ◆ Actually more biologically plausible than sigmoid
 - ◆ **Problem:**
 - ❑ **Not zero-centered output**



Activation Functions:

❖ Leaky ReLU

- ◆ Does not saturate
- ◆ Computationally efficient
- ◆ Converges much faster than sigmoid/tanh in practice! (e.g. 6x)
- ◆ Will not “die”.



Leaky ReLU

$$f(x) = \max(0.01x, x)$$

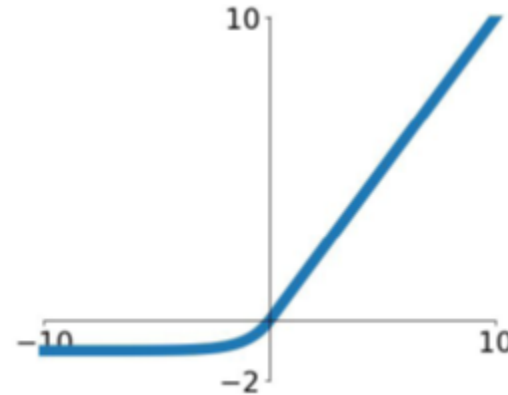
Parametric Rectifier (PReLU)

$$f(x) = \max(\alpha x, x)$$

Activation Functions

❖ Exponential Linear Units (ELU)

- ◆ All benefits of ReLU
- ◆ Closer to zero mean outputs
- ◆ Negative saturation regime compared with Leaky ReLU
- ◆ Adds some robustness to noise



Exponential Linear Units (ELU)

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$

Activation Functions:

❖ **Maxout** : $\max(w_1^T x + b_1, w_2^T x + b_2)$

- ◆ Does not have the basic form of dot product -> nonlinearity
- ◆ Generalizes ReLU and Leaky ReLU
- ◆ Linear Regime! Does not saturate! Does not die!
- ◆ **Problem:**
 - ▣ doubles the number of parameters/neuron

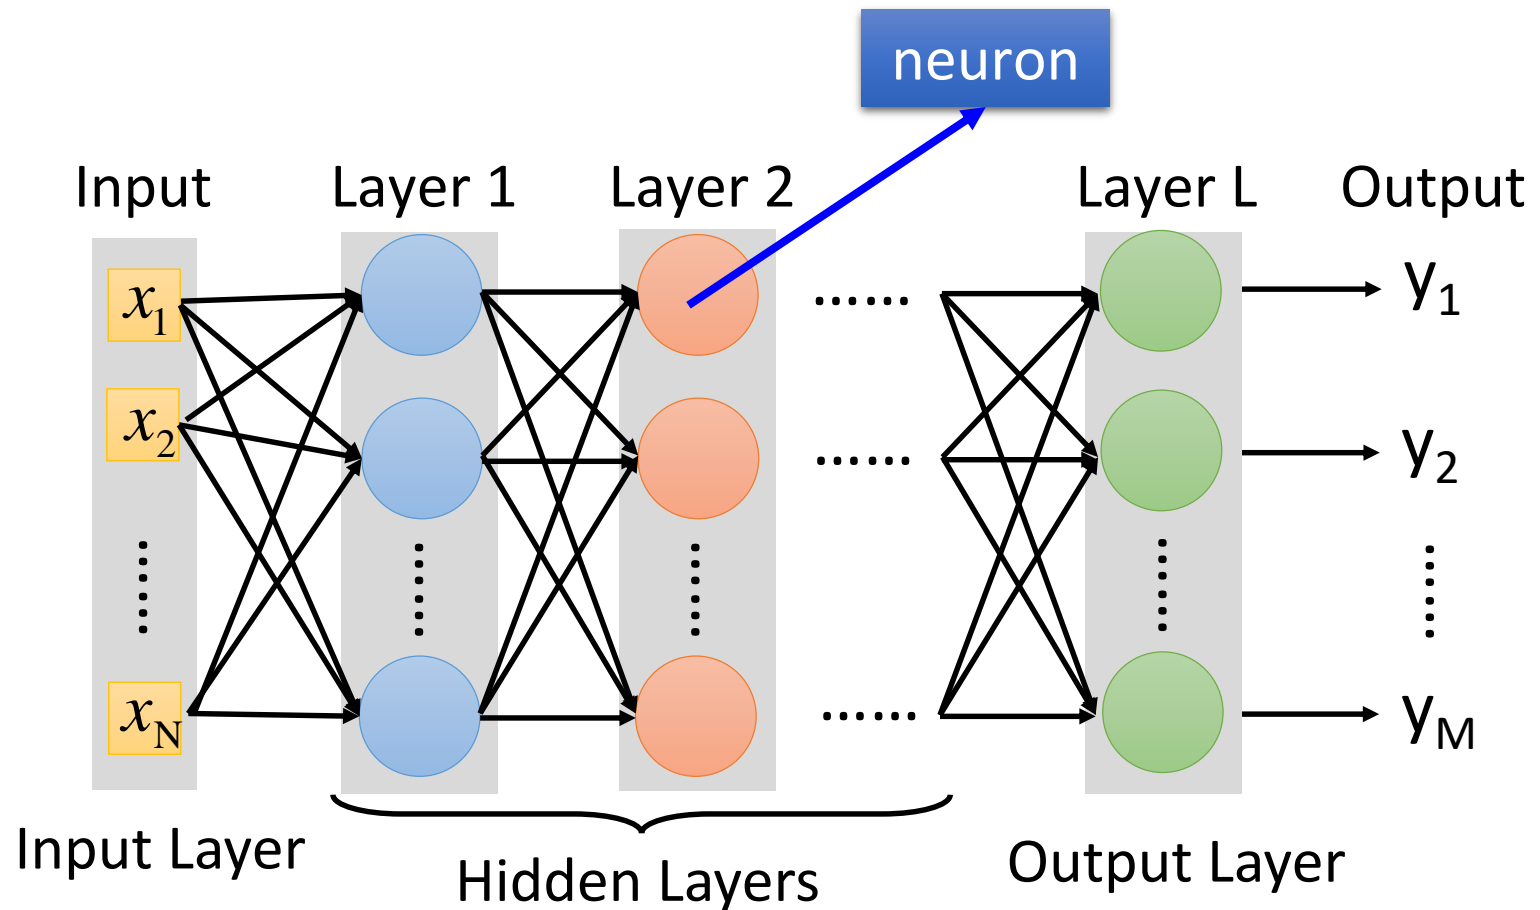
Choosing the right Activation Function

- ❖ Good or bad – there is no rule of thumb
- ❖ Depending upon the properties of the problem, we might be able to make a better choice for **easy and quicker** convergence of the network
 - ◆ Sigmoid functions and their combinations generally work better **in the case of classifiers**
 - ◆ Sigmoid and tanh functions: sometimes to be avoided due to the vanishing gradient problem
 - ◆ ReLU function : a general activation function, used in most cases these days, but only be used in the hidden layers. Be careful with the learning rates
 - ◆ In the case of dead neurons in networks, try out the Leaky ReLU / Maxout / ELU
- ❖ Usually can begin with using ReLU function and then move over to other activation functions in case ReLU doesn't provide with optimum results

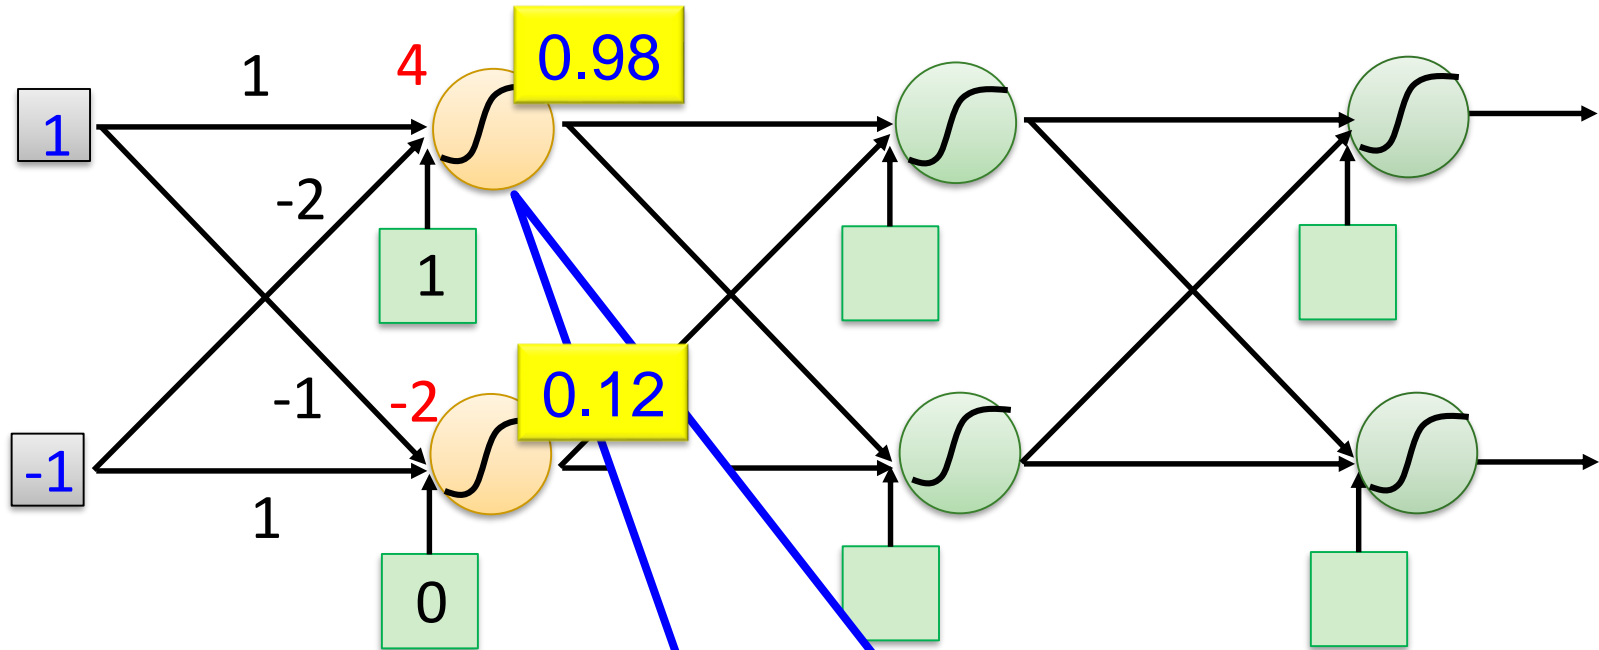
Outline

- ❖ Modeling one neuron
 - ❖ Activation functions
 - ❖ Fully connected feed-forward network
 - ❖ How to train a multi-layer network
 - ❖ Representational power of NN
-

Fully connected Feedforward Network



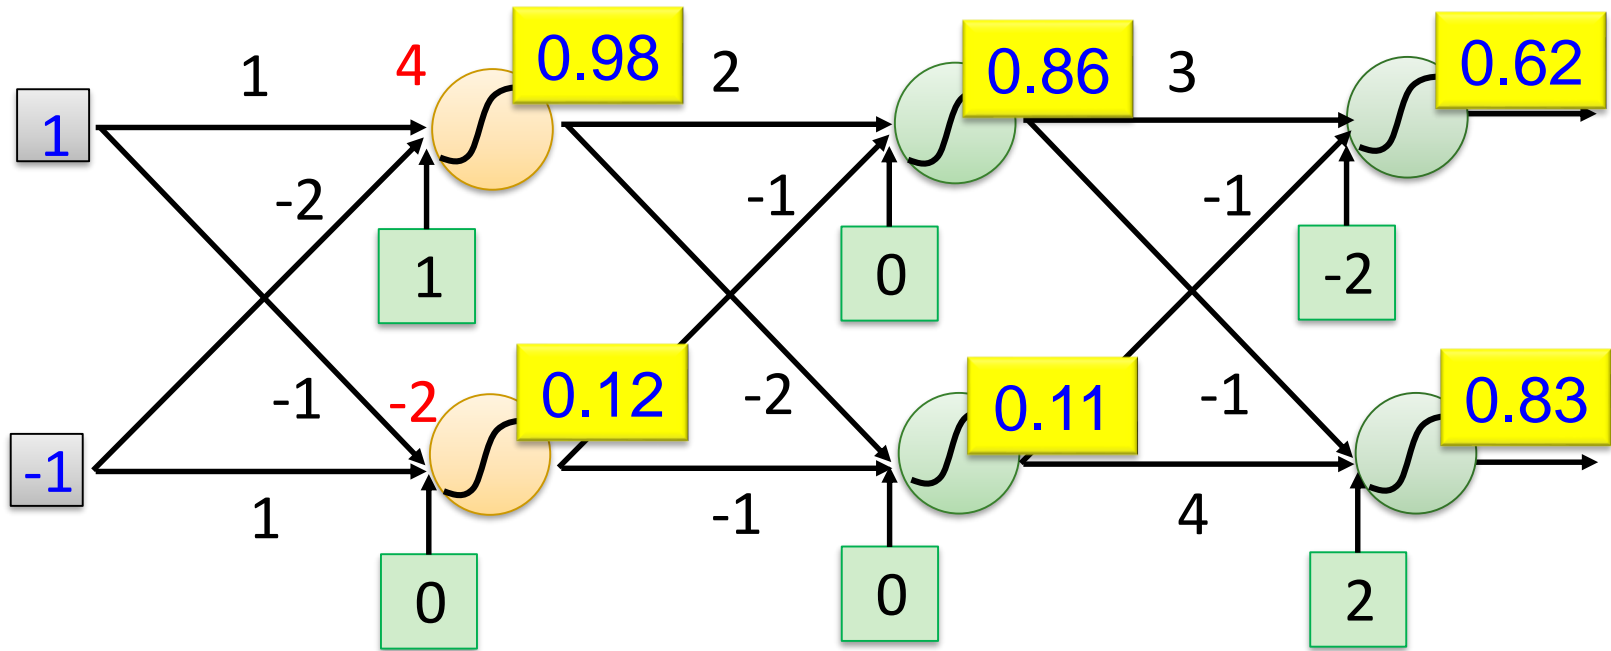
Fully connected Feedforward Network



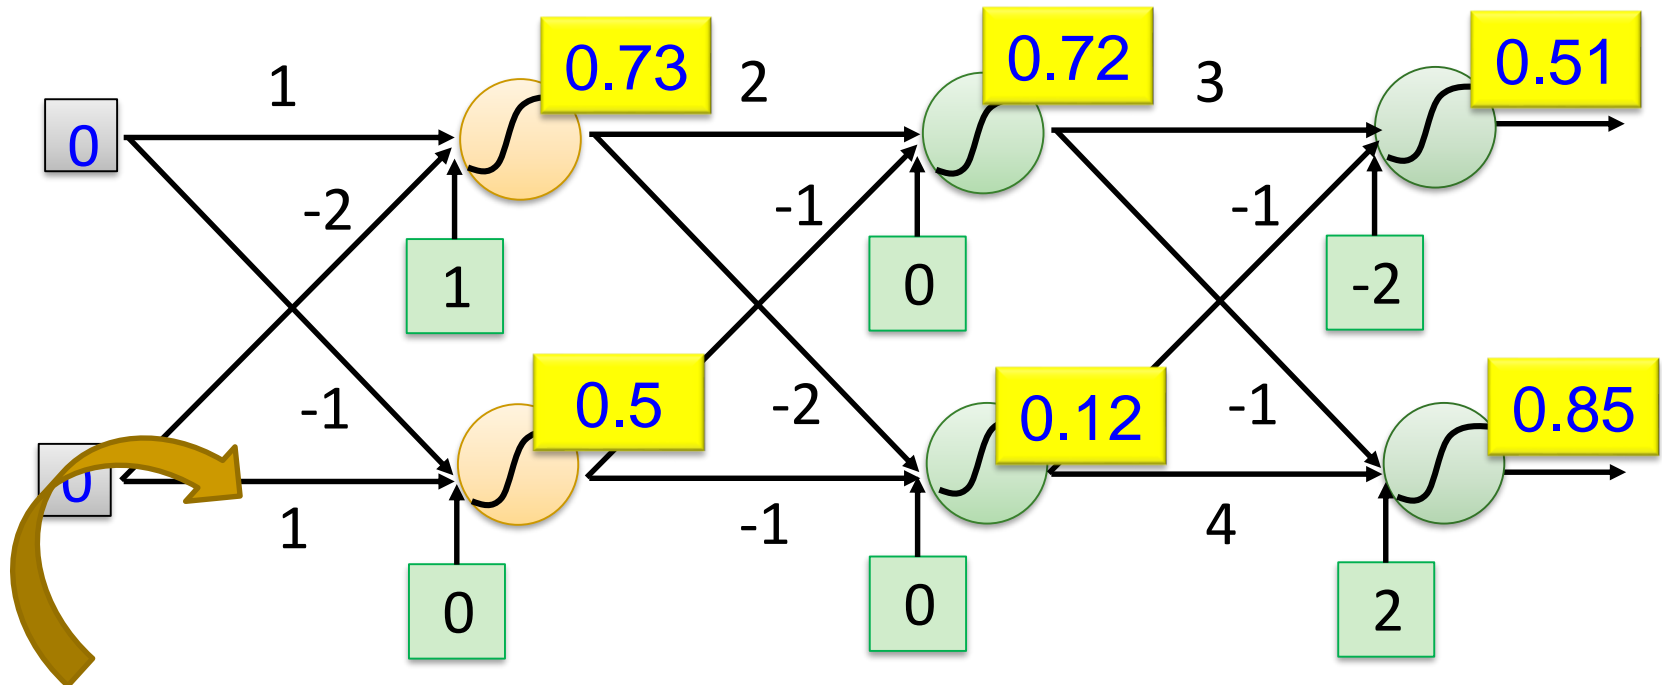
Sigmoid Function $\sigma(z)$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Fully connected Feedforward Network



Fully connected Feedforward Network



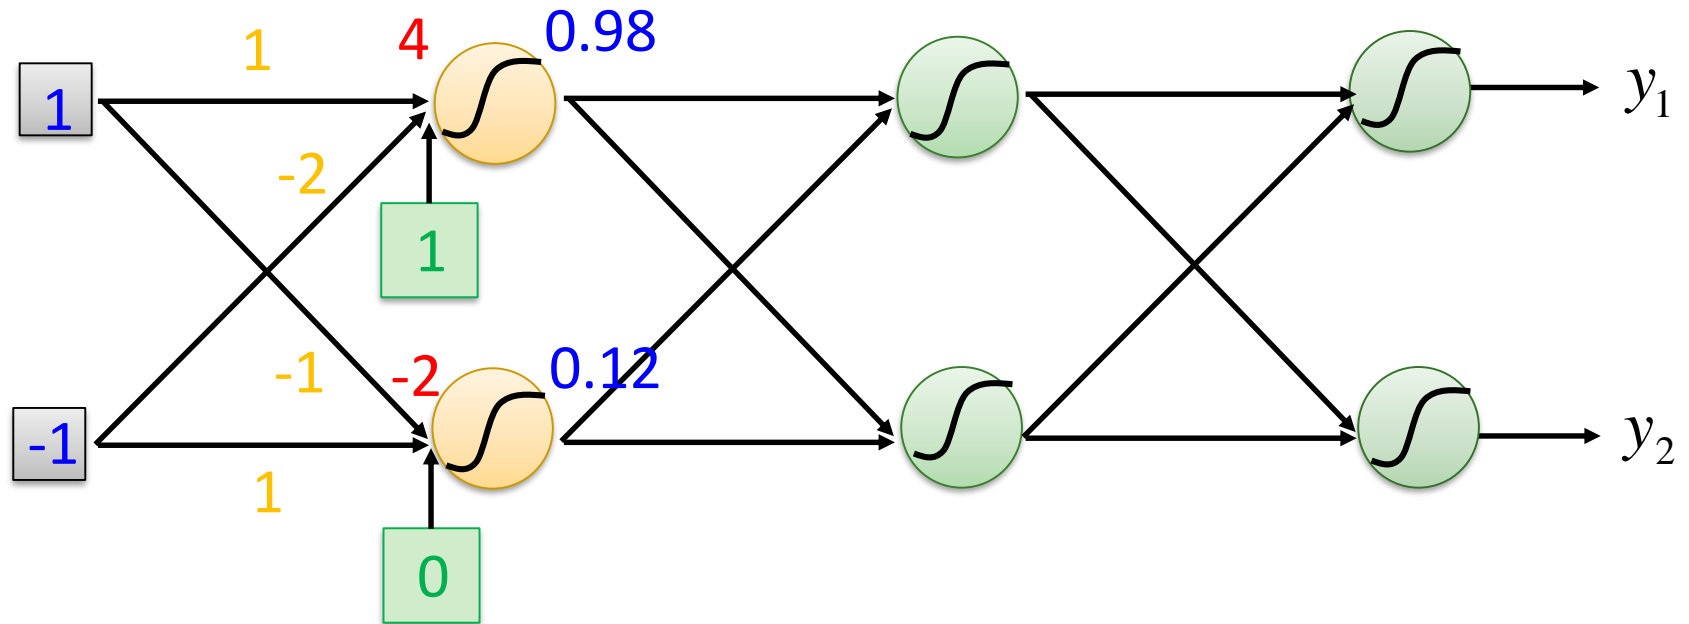
This is a function.

Input vector, output vector

$$f\left(\begin{bmatrix} 1 \\ -1 \end{bmatrix}\right) = \begin{bmatrix} 0.62 \\ 0.83 \end{bmatrix} \quad f\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}\right) = \begin{bmatrix} 0.51 \\ 0.85 \end{bmatrix}$$

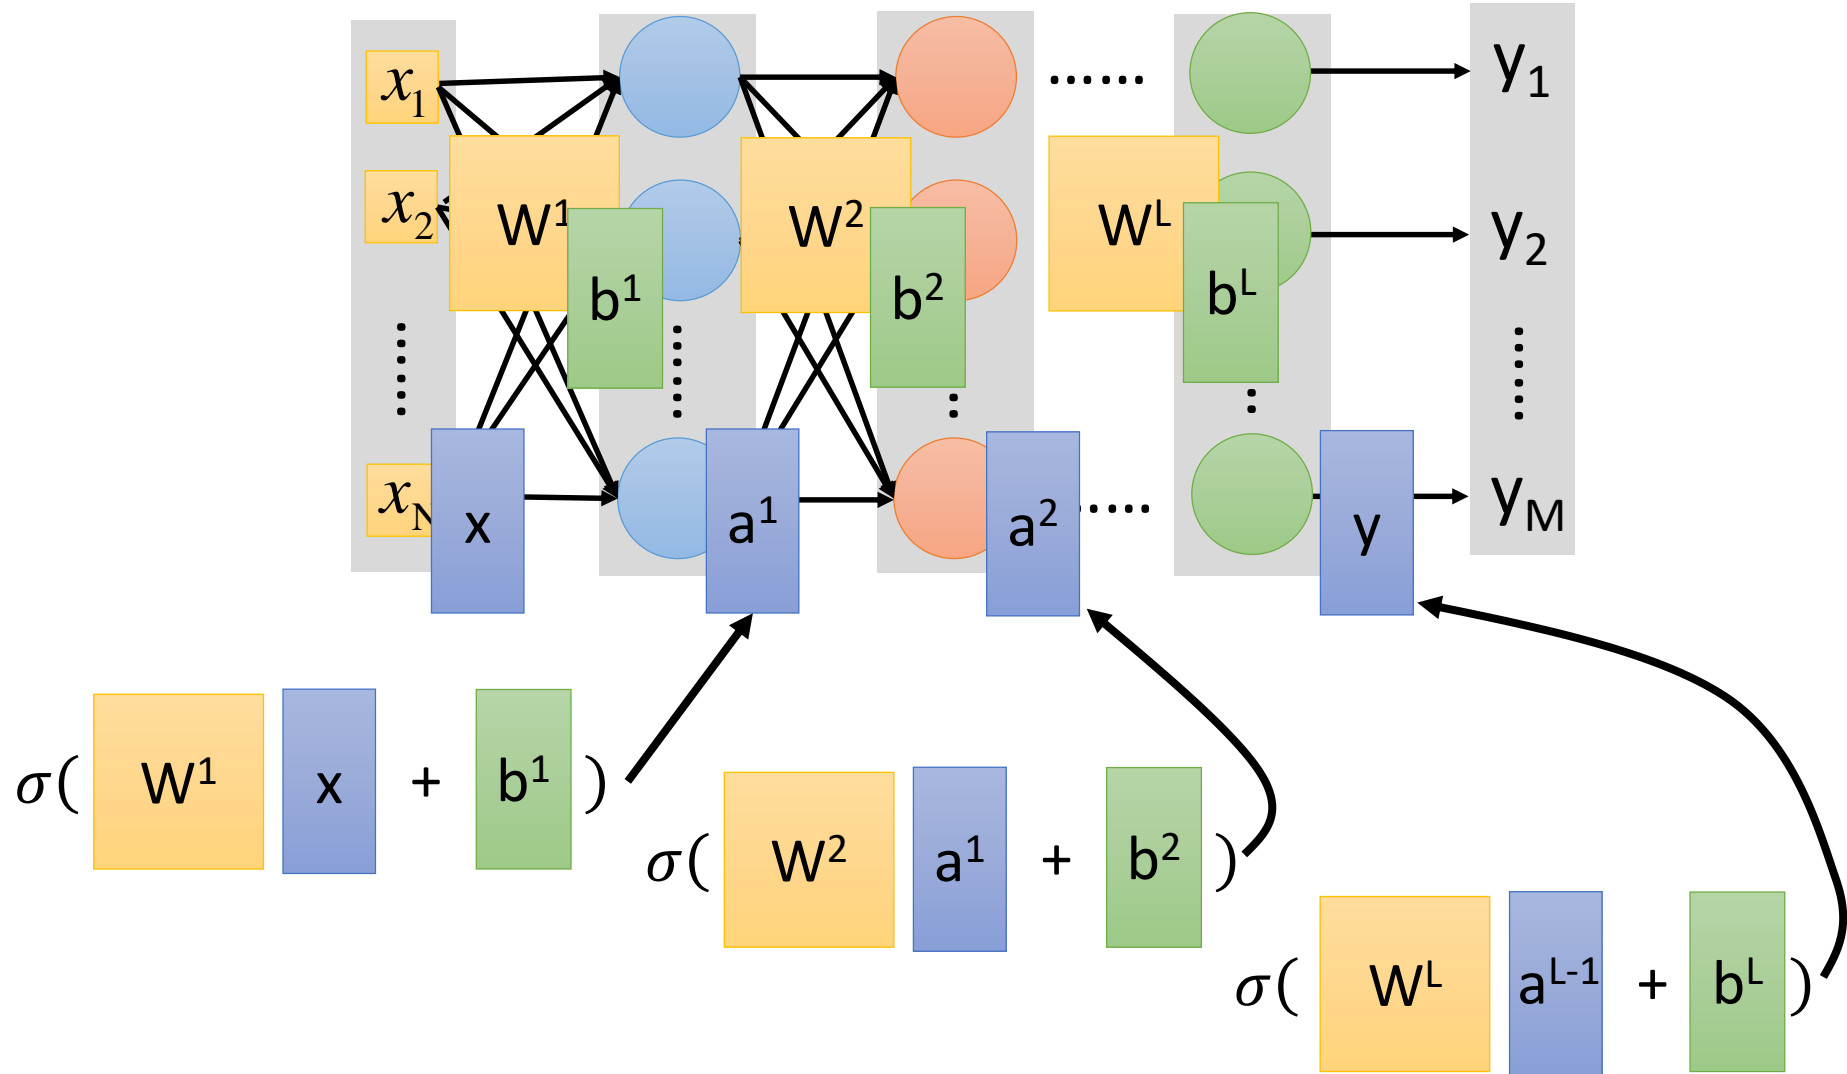
Given network structure, we define a function set

Matrix Operation

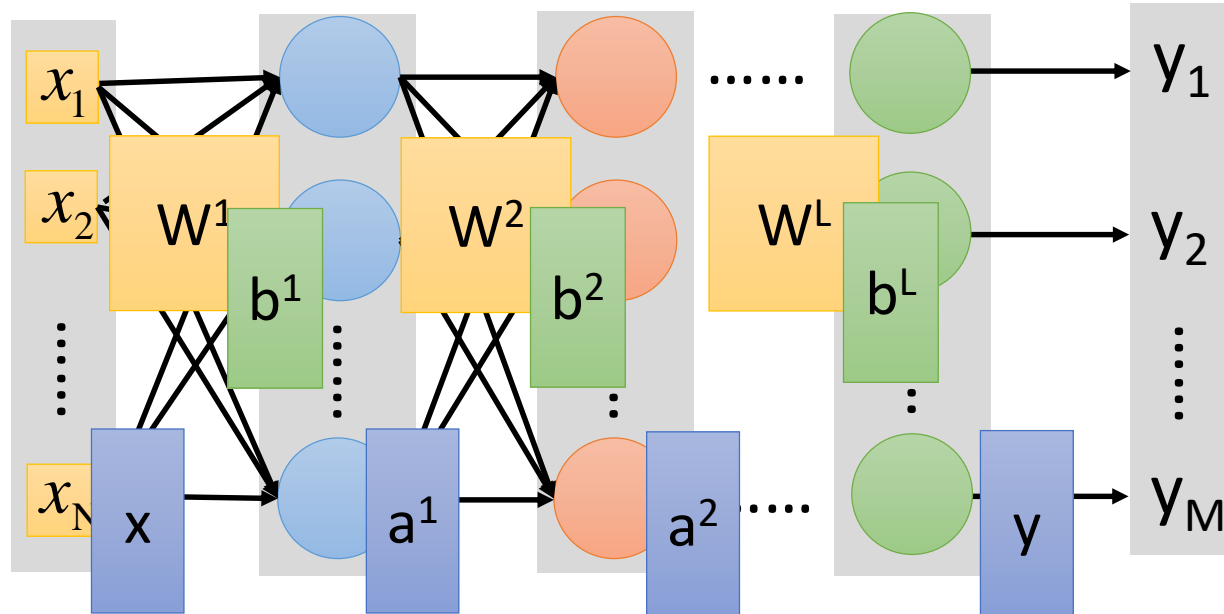


$$\sigma\left(\underbrace{\begin{bmatrix} 1 & -2 \\ -1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix}}_{\begin{bmatrix} 4 \\ -2 \end{bmatrix}} \right) = \begin{bmatrix} 0.98 \\ 0.12 \end{bmatrix}$$

Neural Network



Neural Network



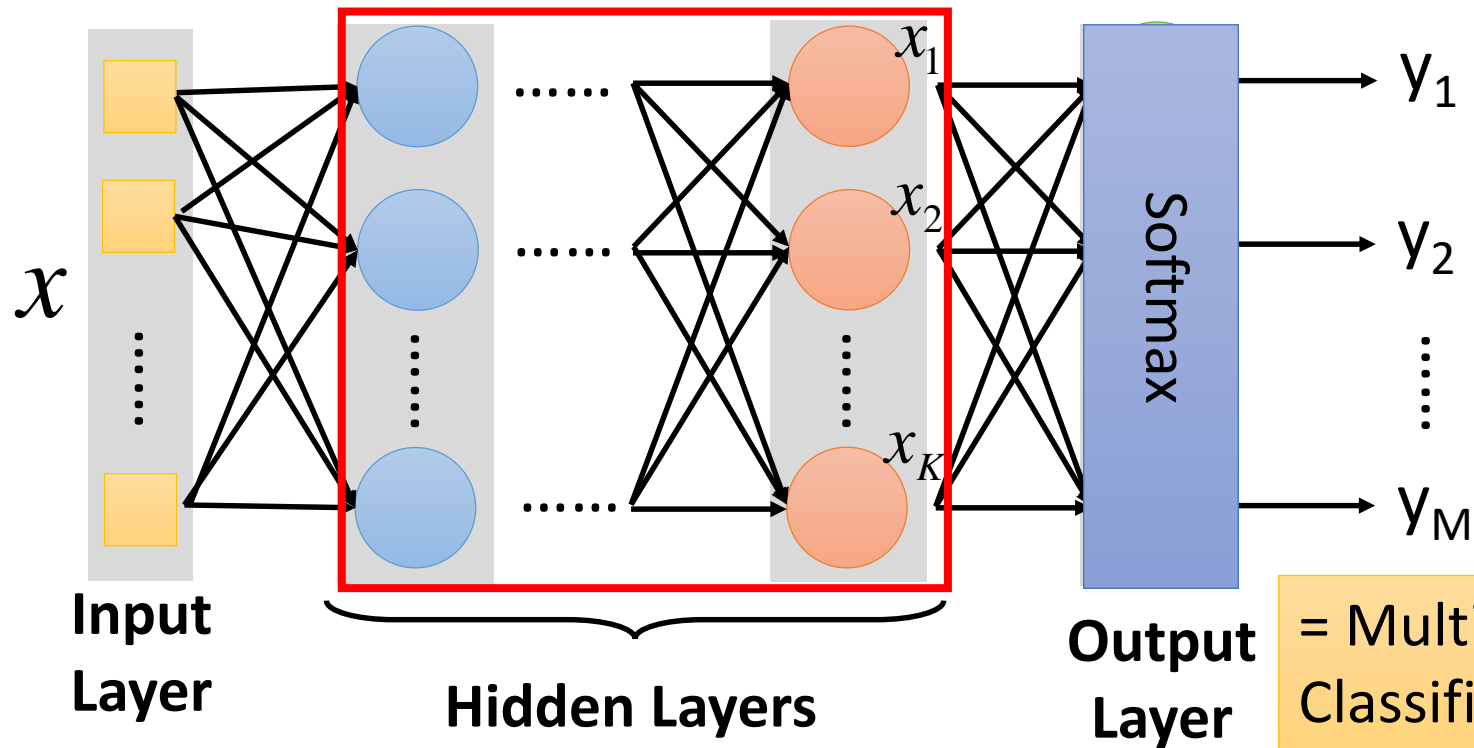
$$y = f(x)$$

Using parallel computing techniques to speed up matrix operation

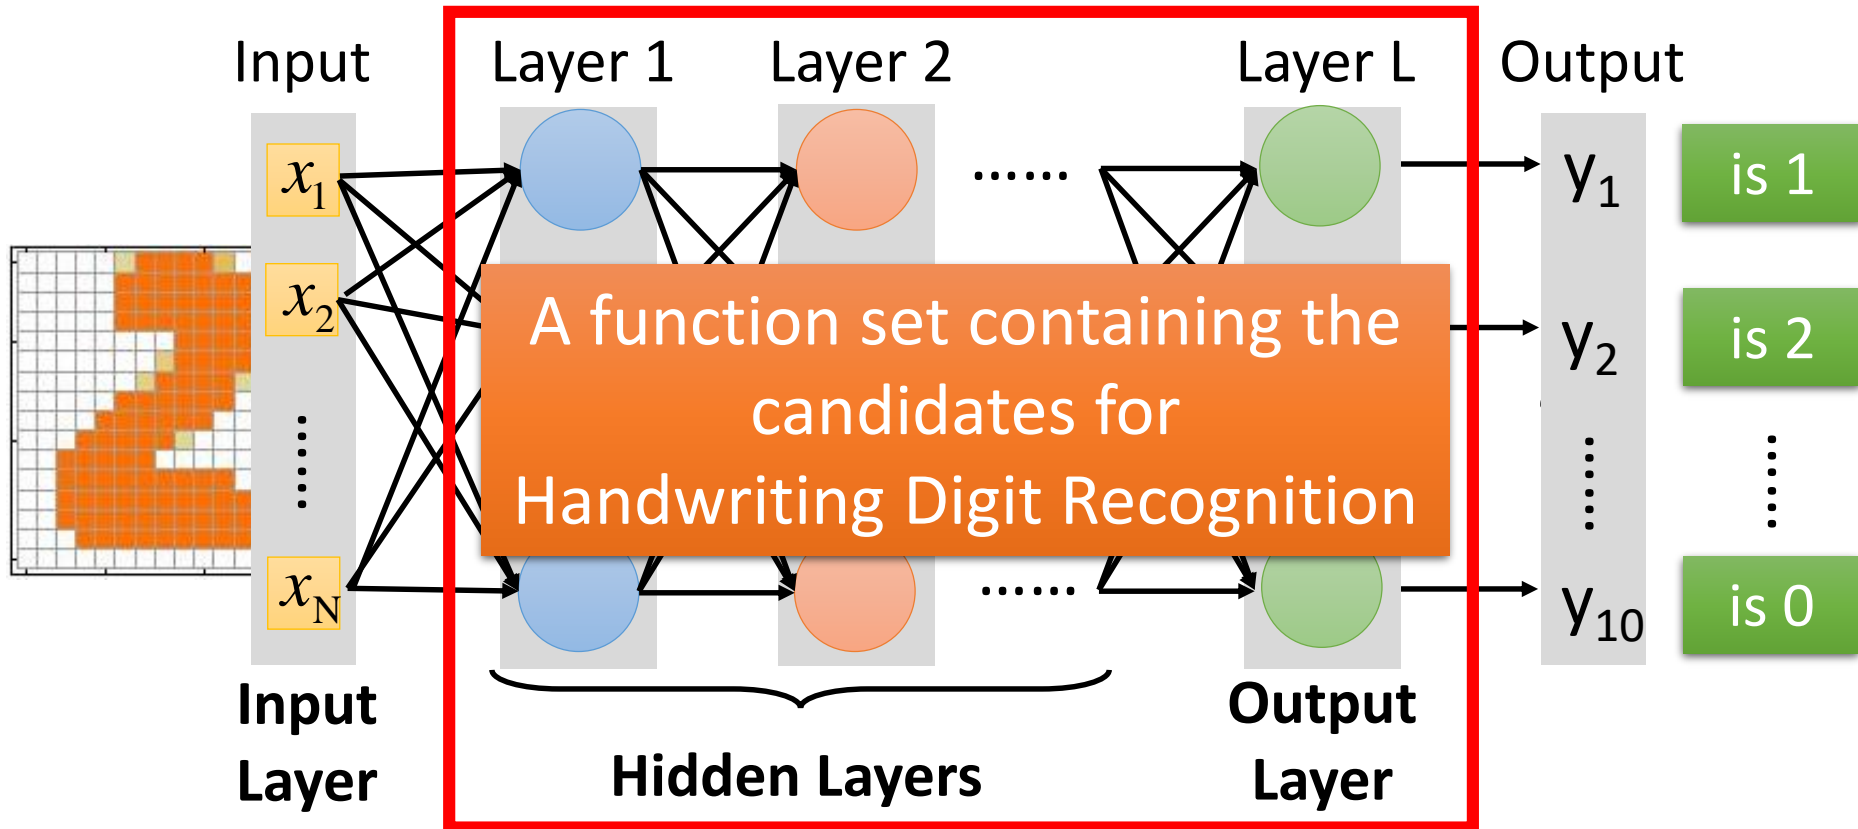
$$= \sigma(W^L \dots \sigma(W^2 \sigma(W^1 x + b^1) + b^2) \dots + b^L)$$

Output Layer as a Multi-Class Classifier

Feature extractor replacing
feature engineering



Example Application



Need to decide the network structure to work well on your dataset.

“Deep” pipeline

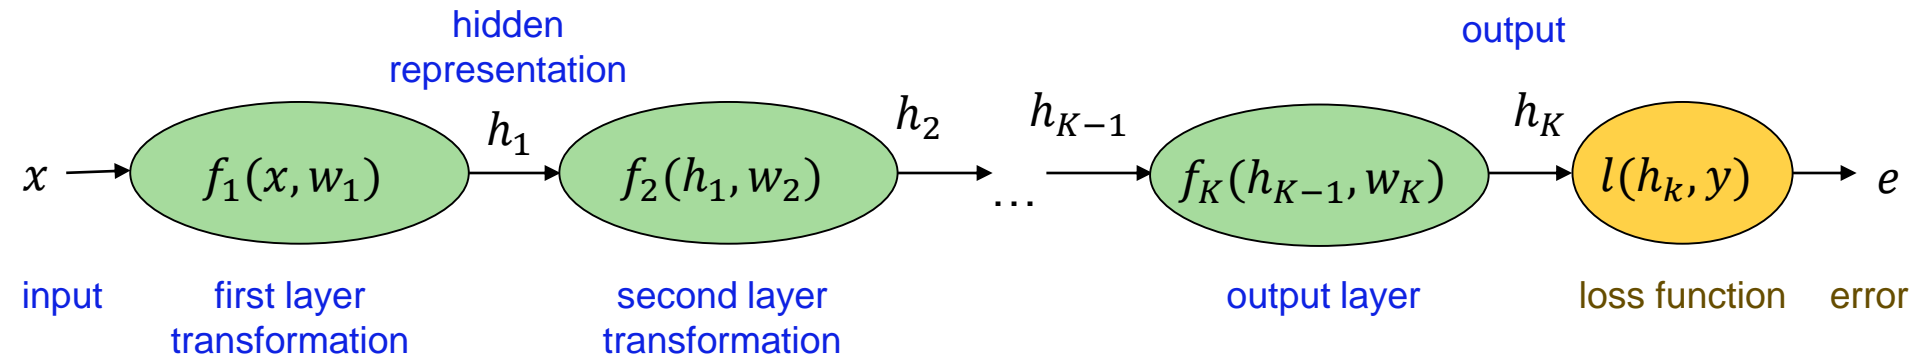


- Learn a *feature hierarchy*
- Each layer extracts features from the output of previous layer
- All layers are trained jointly

Outline

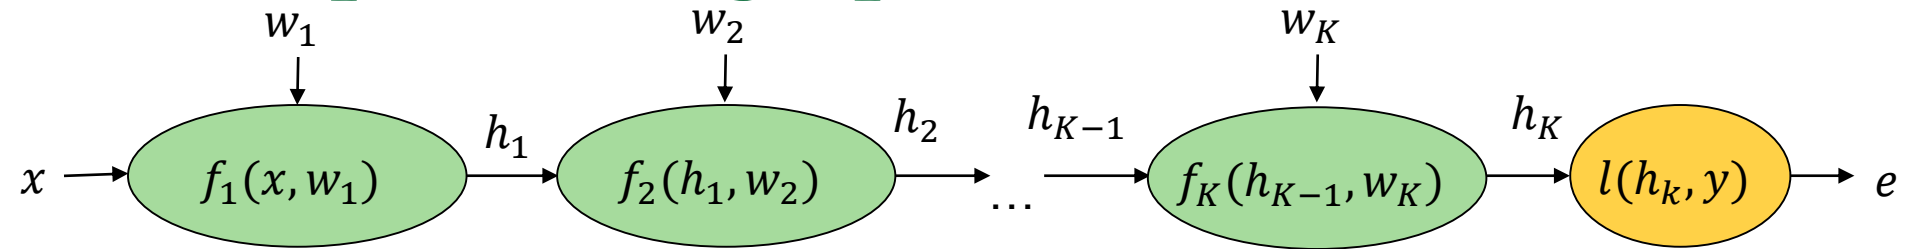
- ❖ Modeling one neuron
 - ❖ Activation functions
 - ❖ Fully connected feed-forward network
 - ❖ How to train a multi-layer network
 - ❖ Representational power of NN
-

How to train a multi-layer network?



- ❖ We need to find the gradient of the error w. r.t. the parameters of each layer, $\frac{\partial e}{\partial w_k}$, to perform updates $w_k \leftarrow w_k - \eta \frac{\partial e}{\partial w_k}$

Computation graph



Chain Rule

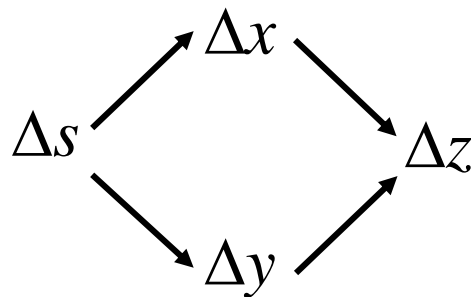
Case 1 $y = g(x)$ $z = h(y)$

$$\Delta x \rightarrow \Delta y \rightarrow \Delta z$$

$$\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$$

Case 2

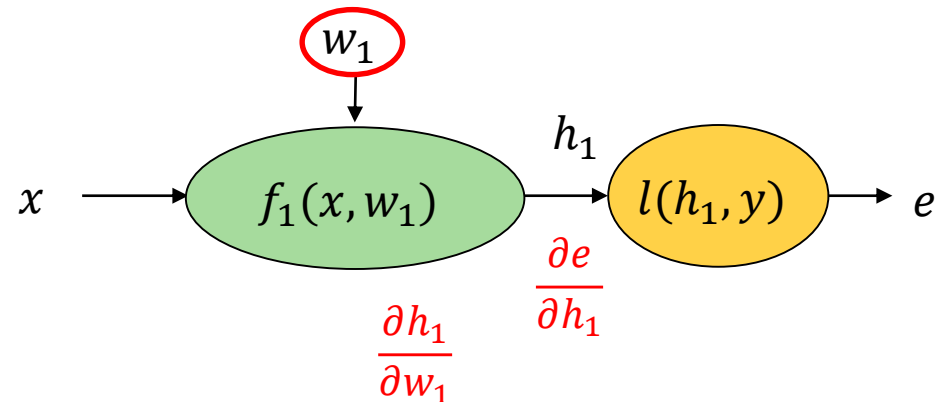
$$x = g(s) \quad y = h(s) \quad z = k(x, y)$$



$$\frac{dz}{ds} = \frac{\partial z}{\partial x} \frac{dx}{ds} + \frac{\partial z}{\partial y} \frac{dy}{ds}$$

Chain rule

Let's start with $k = 1$



❖ $e = l(f_1(x, w_1), y)$

❖ $\frac{\partial}{\partial w_1} l(f_1(x, w_1), y) = \frac{\partial e}{\partial h_1} \frac{\partial h_1}{\partial w_1}$

❖ Example: $e = (y - w_1^T x)^2$

❖ $h_1 = f_1(x, w_1) = w_1^T x$

❖ $e = l(h_1, y) = (y - h_1)^2$

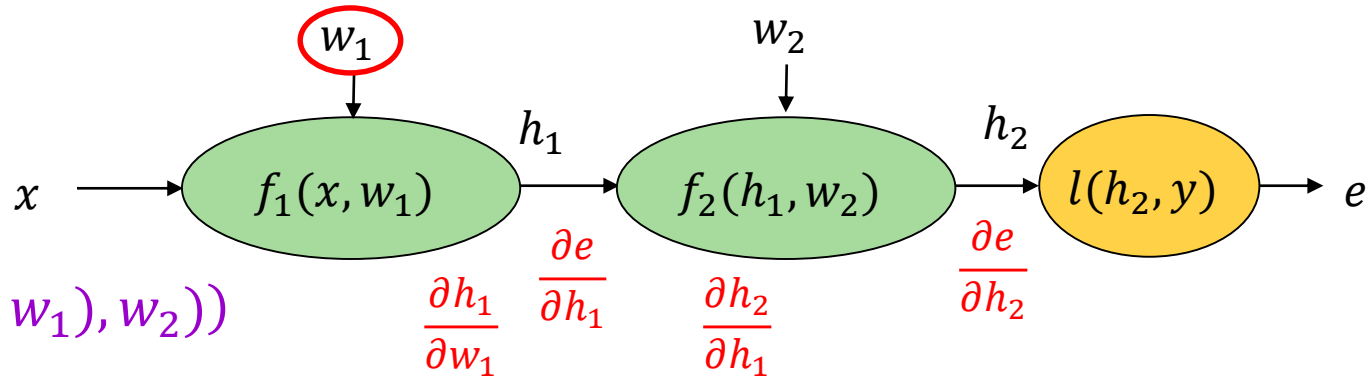
$$\frac{\partial h_1}{\partial w_1} = x$$

$$\frac{\partial e}{\partial h_1} = -2(y - h_1) = -2(y - w_1^T x)$$

❖ $\frac{\partial e}{\partial w_1} = \frac{\partial e}{\partial h_1} \frac{\partial h_1}{\partial w_1} = -2x(y - w_1^T x)$

Chain rule

$k = 2$



❖ $e = l(f_2(f_1(x, w_1), w_2))$

$$\frac{\partial e}{\partial w_2} = \frac{\partial e}{\partial h_2} \frac{\partial h_2}{\partial w_2}$$

$$\frac{\partial e}{\partial w_1} = \frac{\partial e}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial w_1}$$

❖ Example: $e = -\log(\sigma(w_1^T x))$ (assume $y = 1$)

$$h_1 = f_1(x, w_1) = w_1^T x$$

$$h_2 = f_2(h_1, w_2) = \sigma(h_1)$$

$$e = l(h_2, 1) = -\log(h_2)$$

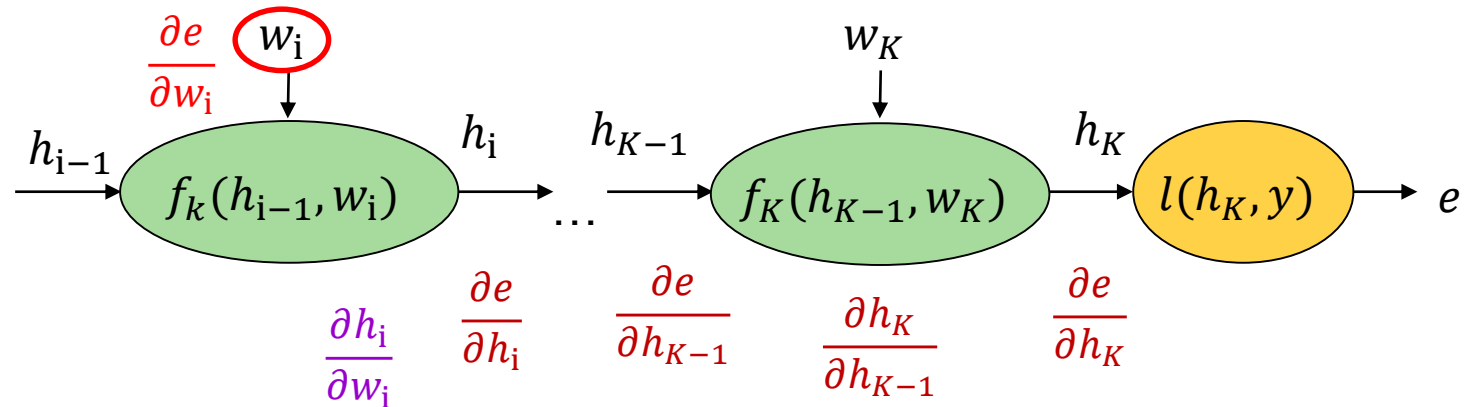
$$\frac{\partial h_1}{\partial w_1} = x$$

$$\frac{\partial h_2}{\partial h_1} = \sigma'(h_1) = \sigma(h_1)(1 - \sigma(h_1))$$

$$\frac{\partial e}{\partial h_2} = -\frac{1}{h_2}$$

$$\frac{\partial e}{\partial w_1} = \frac{\partial e}{\partial h_2} \frac{\partial h_2}{\partial h_1} \frac{\partial h_1}{\partial w_1} = -\frac{1}{\sigma(w_1^T x)} \sigma(w_1^T x) (1 - \sigma(w_1^T x)) x = \sigma(-w_1^T x) x - x$$

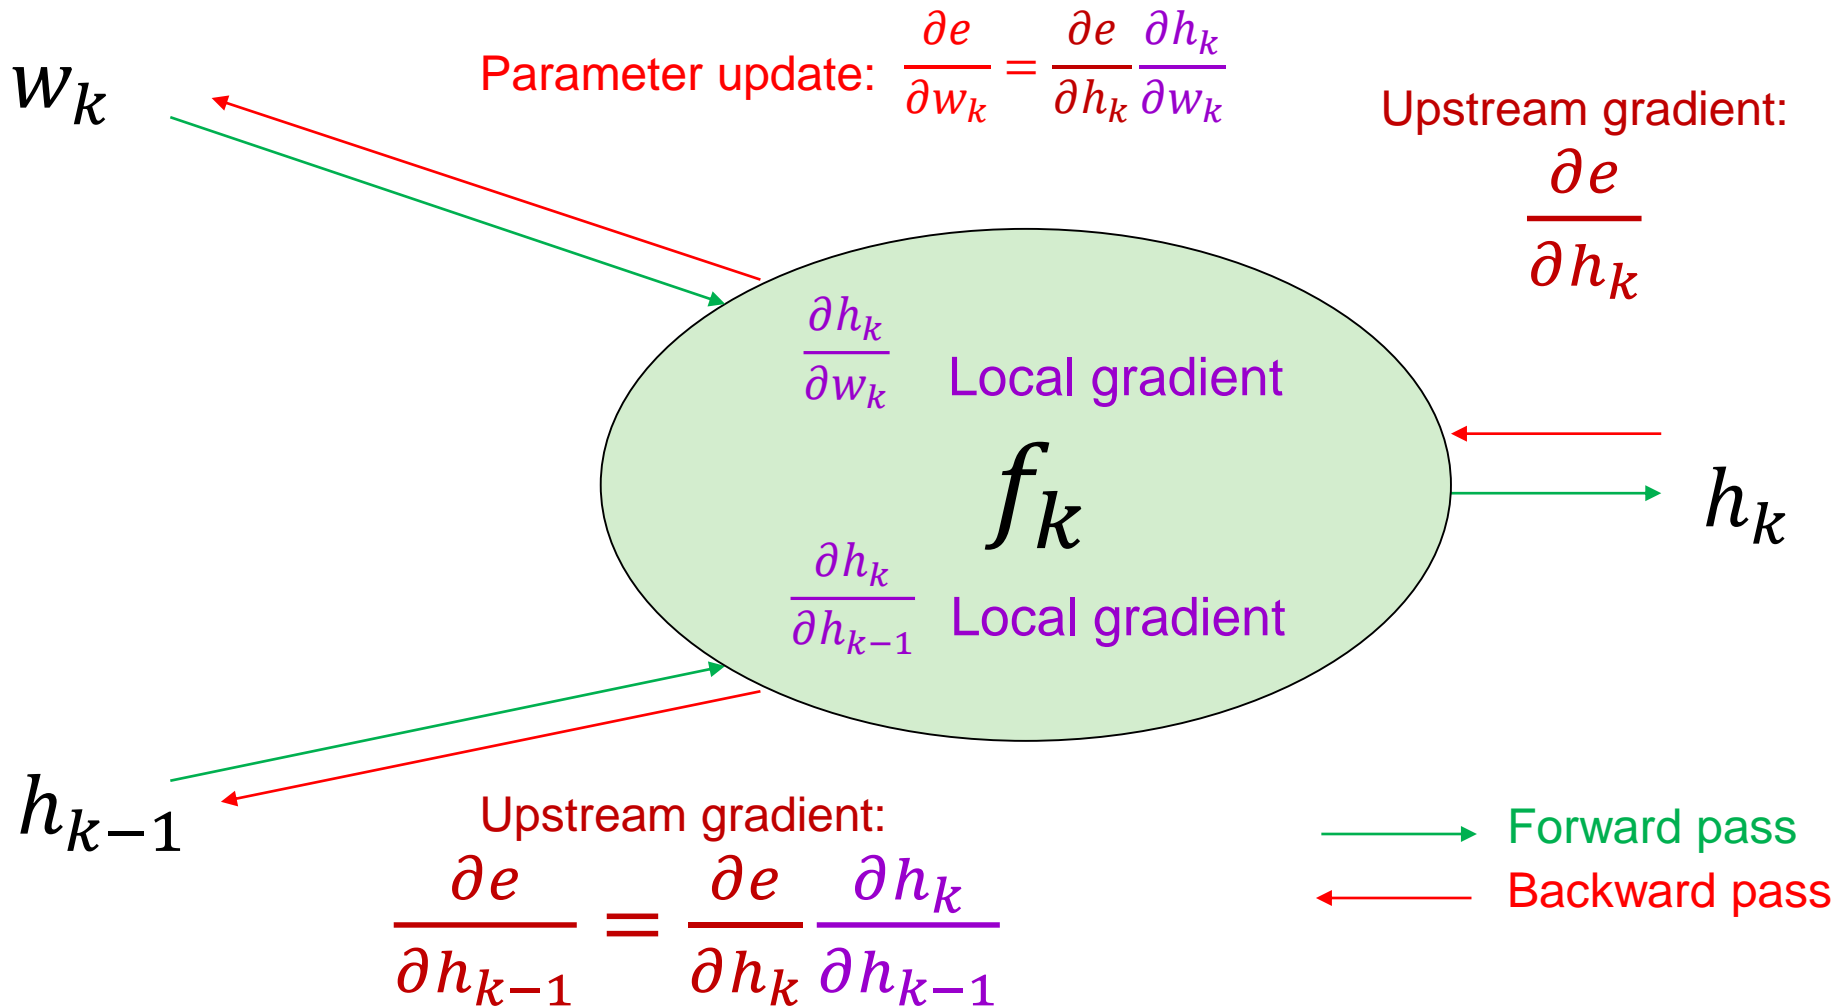
Chain rule



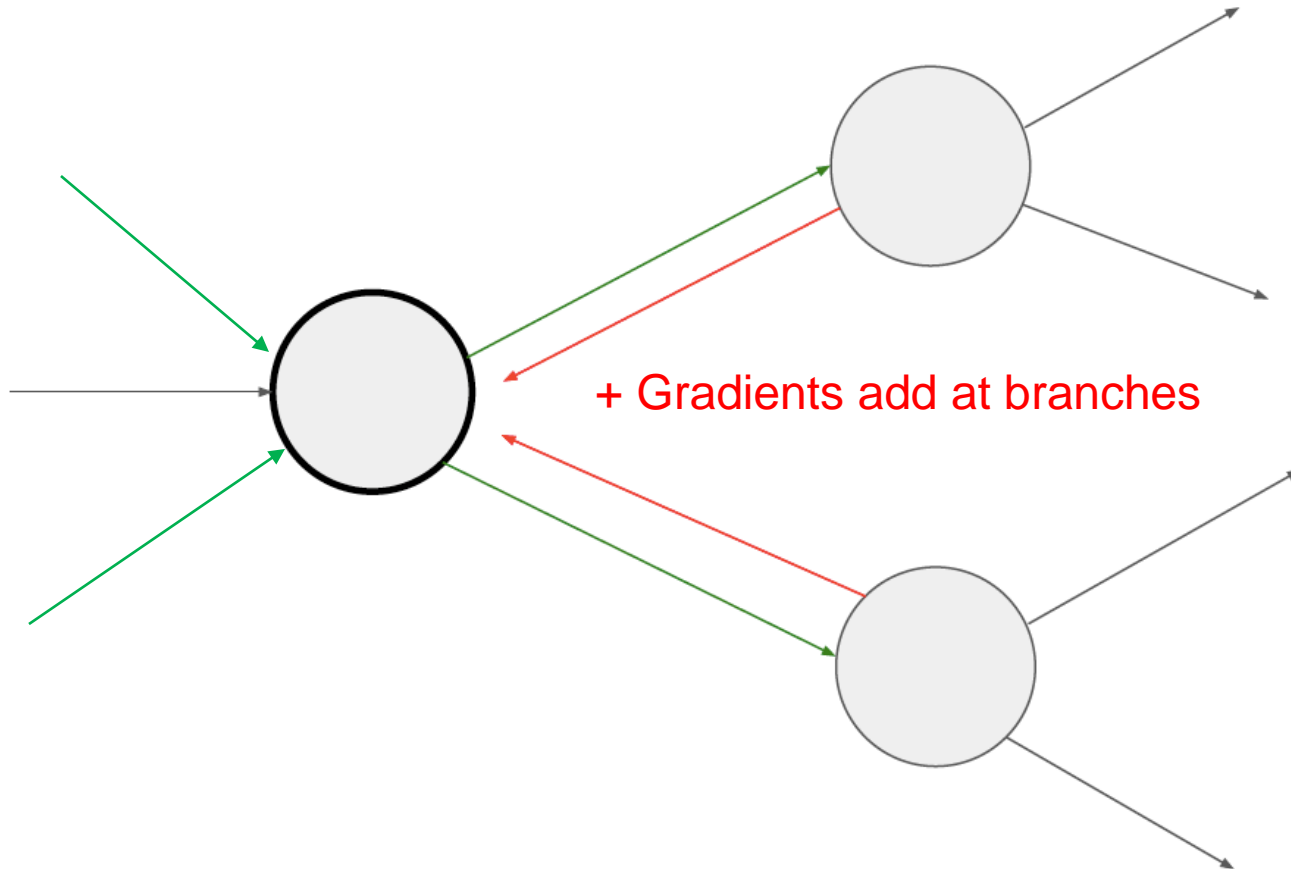
❖ General case:

$$\text{❖ } \frac{\partial e}{\partial w_i} = \underbrace{\frac{\partial e}{\partial h_K} \frac{\partial h_K}{\partial h_{K-1}} \cdots \frac{\partial h_{i+1}}{\partial h_i}}_{\text{Upstream gradient } \frac{\partial e}{\partial h_i}} \underbrace{\frac{\partial h_i}{\partial w_i}}_{\text{Local gradient}}$$

Backpropagation summary

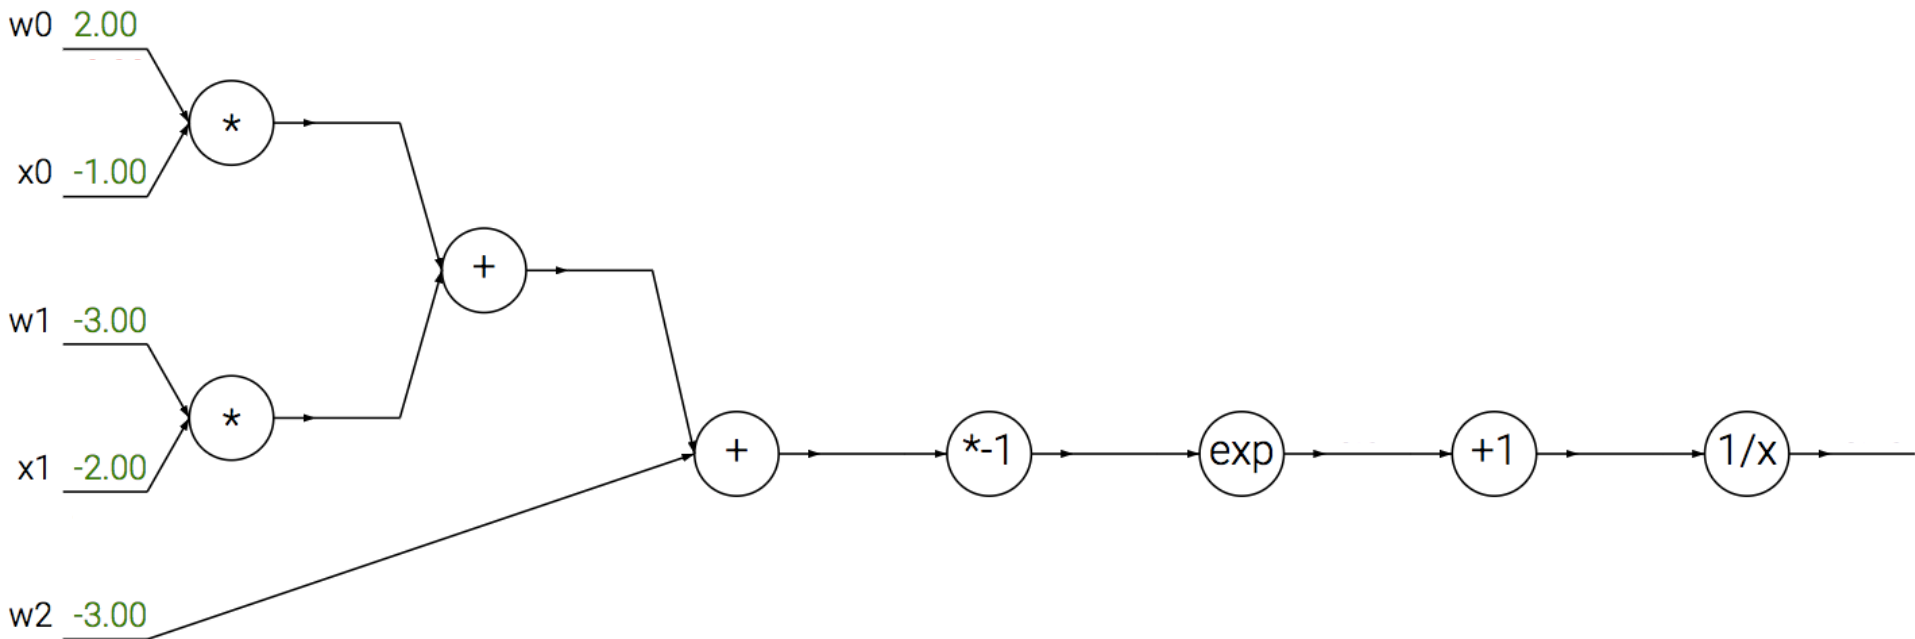


What about more general computation graphs?



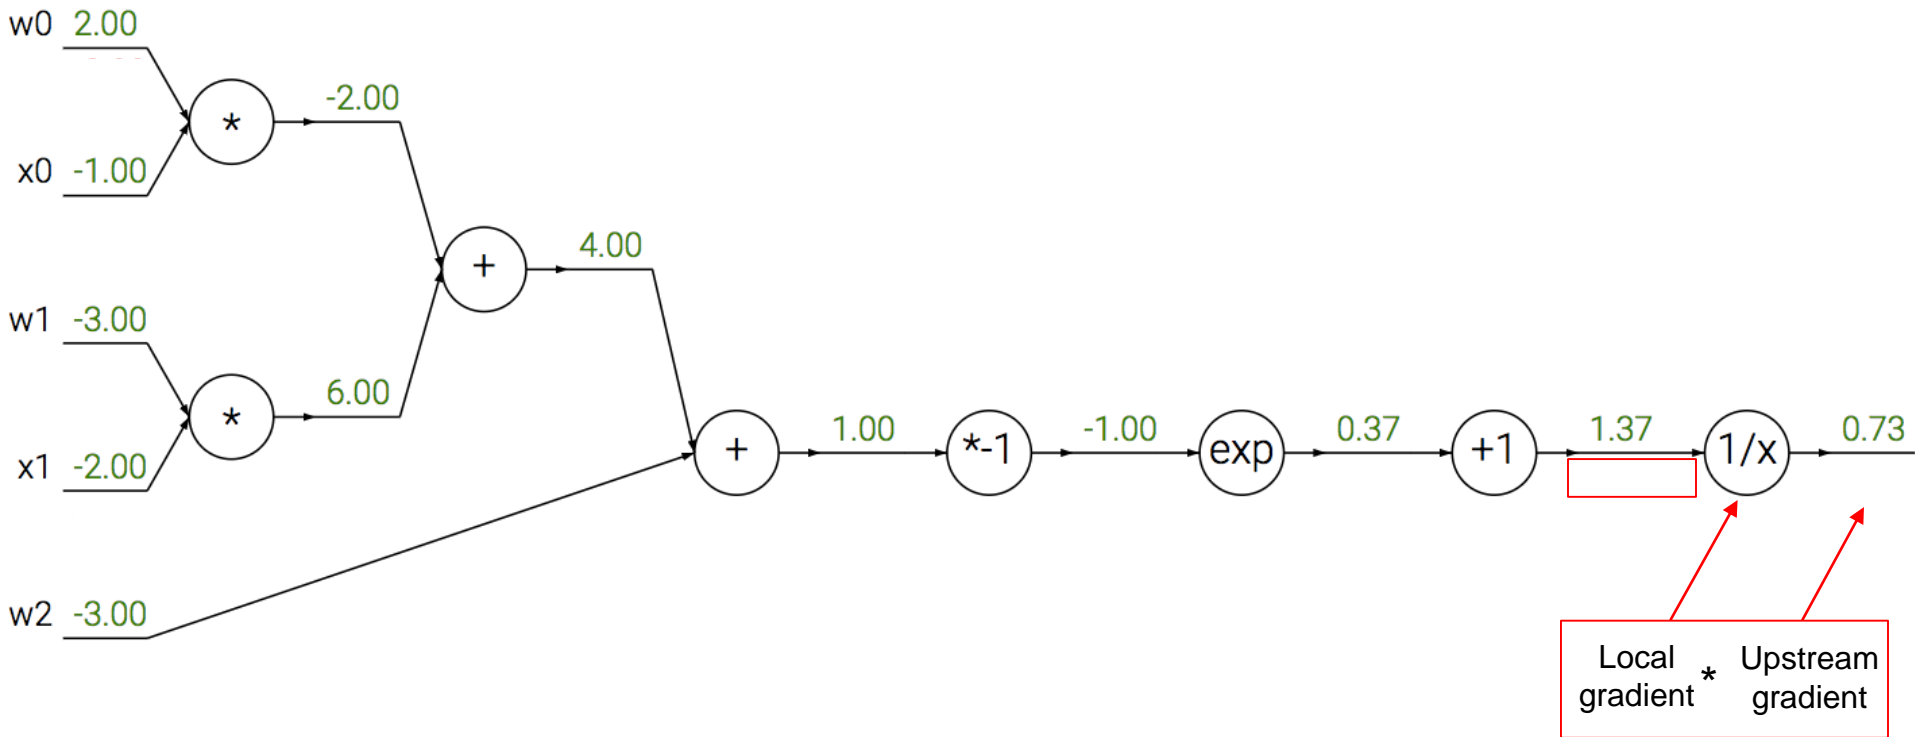
A detailed example

$$f(x, w) = \frac{1}{1 + \exp[-(w_0x_0 + w_1x_1 + w_2)]}$$



A detailed example

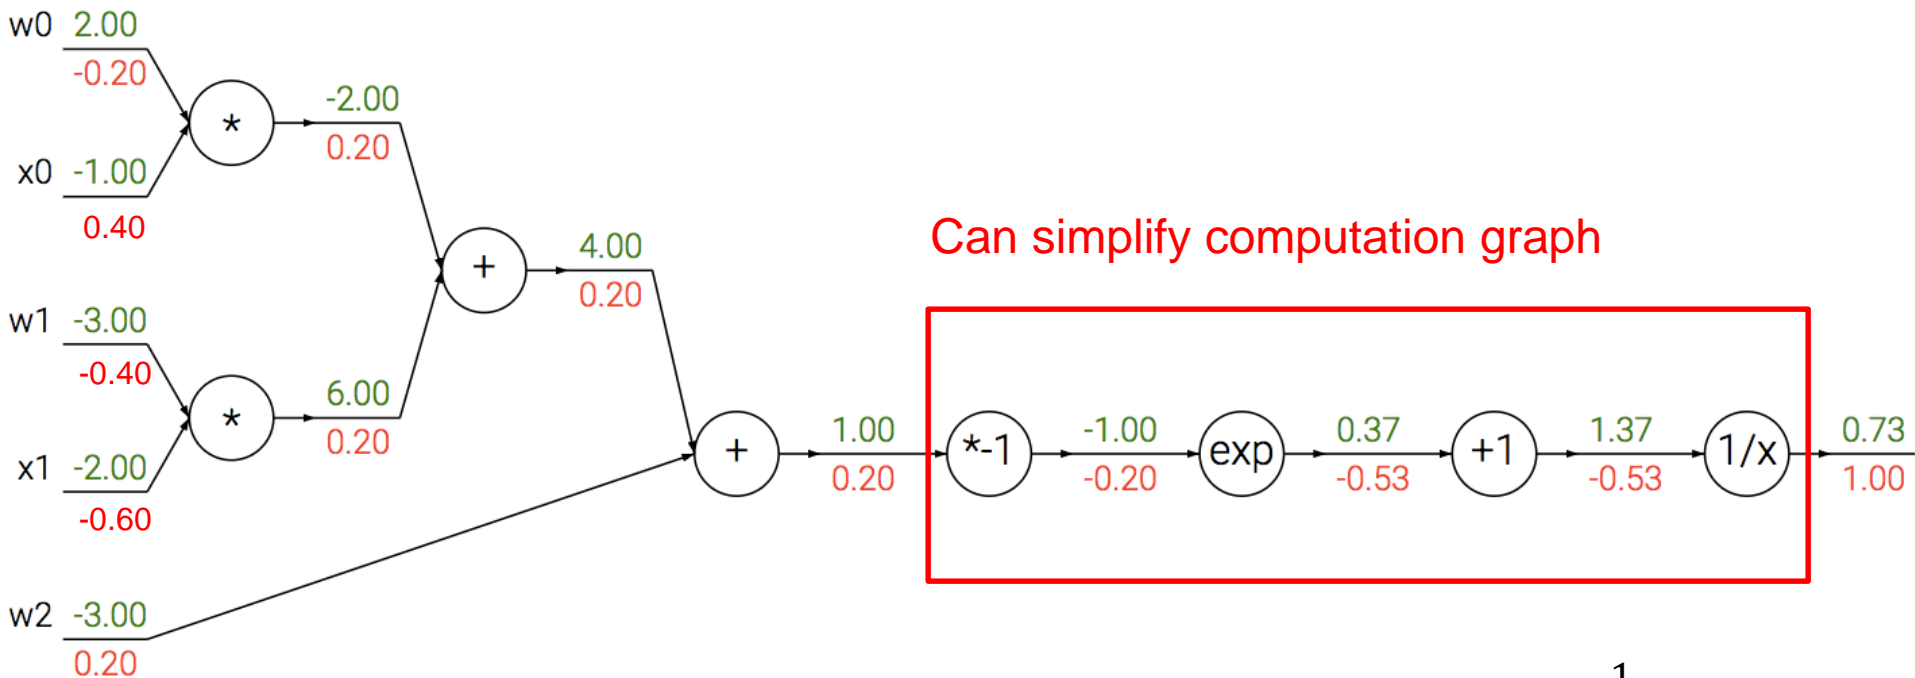
$$f(x, w) = \frac{1}{1 + \exp[-(w_0x_0 + w_1x_1 + w_2)]}$$



$$(1/x)' = -1/x^2$$
$$-\frac{1}{1.37^2} * 1 = -0.53$$

A detailed example

$$f(x, w) = \frac{1}{1 + \exp[-(w_0x_0 + w_1x_1 + w_2)]}$$

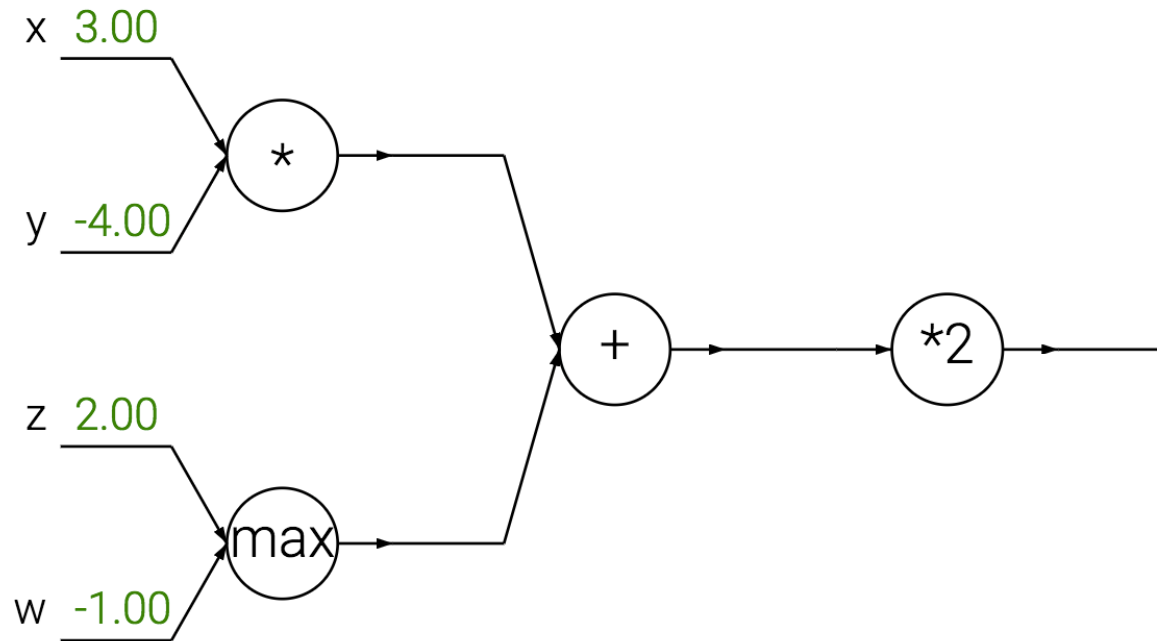


$$\text{Sigmoid gate } \sigma(x) = \frac{1}{1 + \exp(-x)}$$

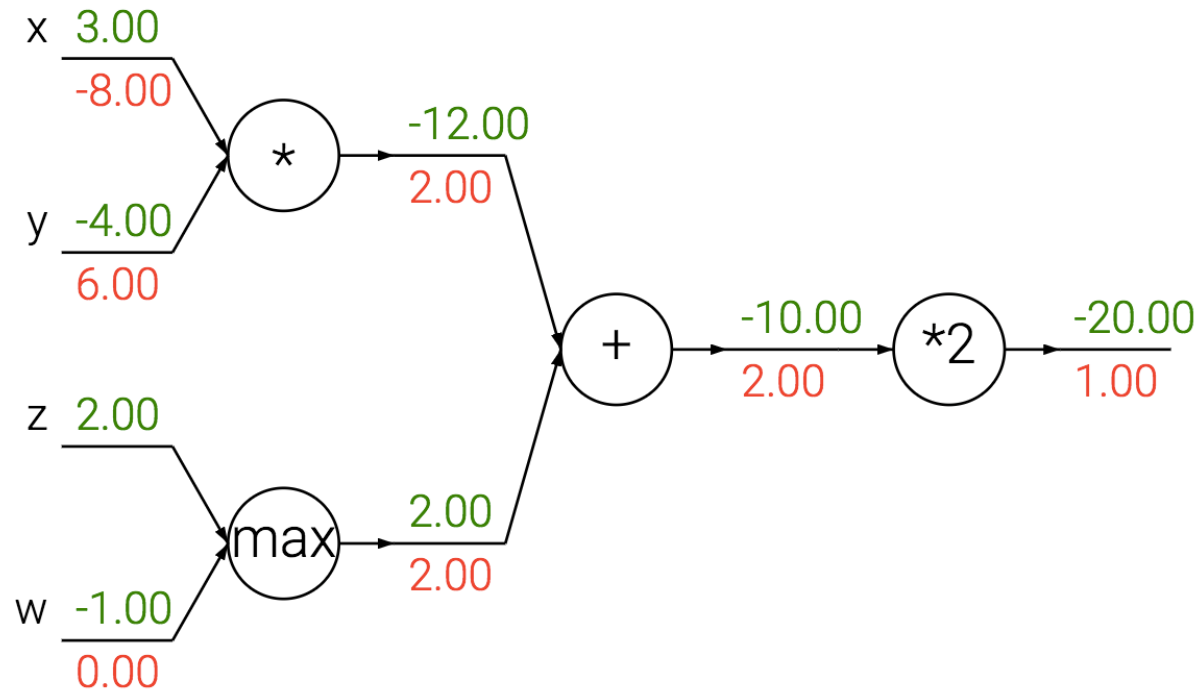
$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

$$\sigma(1)(1 - \sigma(1)) = 0.73 * (1 - 0.73) = 0.20$$

Patterns in gradient flow



Patterns in gradient flow



Add gate: “gradient distributor”

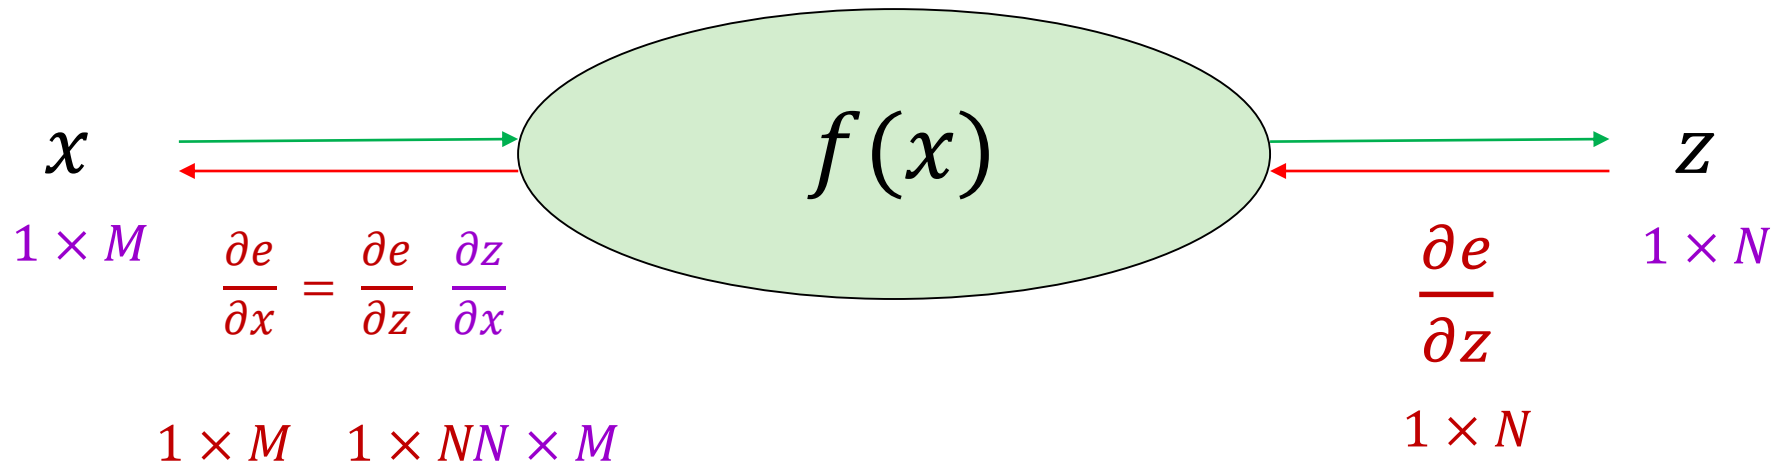
Multiply gate: “gradient switcher”

Max gate: “gradient router”

Dealing with vectors

$$\frac{\partial z}{\partial x} = \begin{pmatrix} \frac{\partial z_1}{\partial x_1} & \cdots & \frac{\partial z_1}{\partial x_M} \\ \vdots & \ddots & \vdots \\ \frac{\partial z_N}{\partial x_1} & \cdots & \frac{\partial z_N}{\partial x_M} \end{pmatrix}$$

$N \times M$
Jacobian



Matrix-vector multiplication

$$\frac{\partial e}{\partial W} = \frac{\partial e}{\partial z} \frac{\partial z}{\partial W}$$

$$M \times N \quad 1 \times N \quad N \times (M \times N) \quad \frac{\partial z}{\partial W}$$

W
 $M \times N$

$N \times (M \times N)$

$$f(x, W) = xW$$

z
 $1 \times N$
 $\frac{\partial e}{\partial z}$

$1 \times N$

$\frac{\partial z}{\partial x}$

$N \times M$

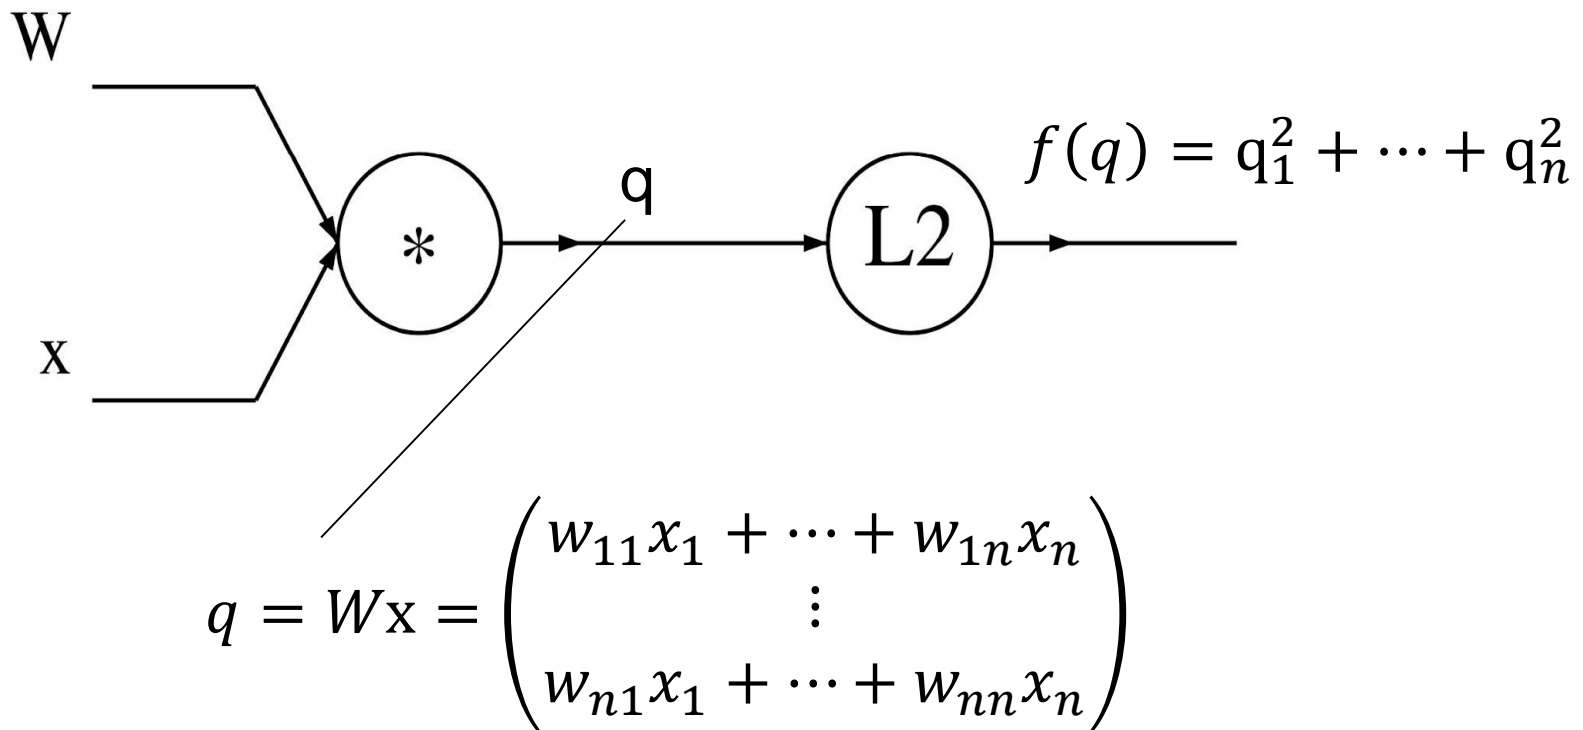
x
 $1 \times M$

$$\frac{\partial e}{\partial x} = \frac{\partial e}{\partial z} \frac{\partial z}{\partial x}$$

$$1 \times M \quad 1 \times N \quad N \times M$$

A vectorized example:

$$f(\mathbf{x}, W) = \sum_{i=1}^n (W \cdot \mathbf{x})_i^2$$

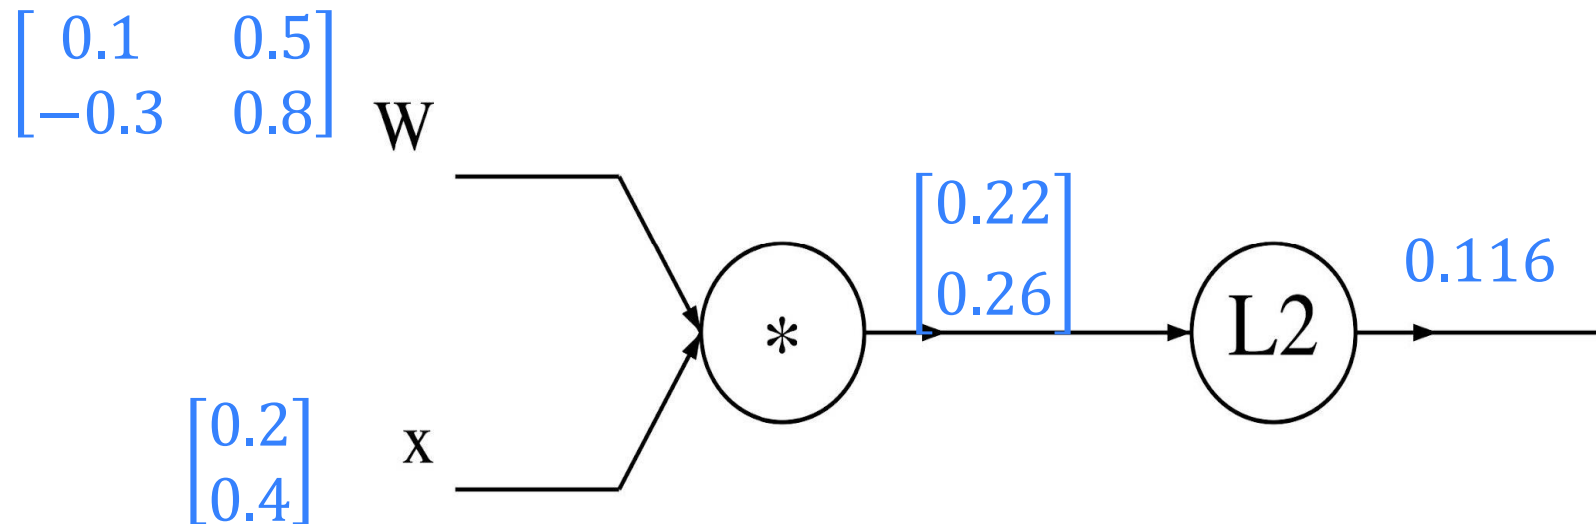


A vectorized example: $f(x, W) = \sum_{i=1}^n (W \cdot x)_i^2$

Feed-forward:

$$q = Wx = \begin{pmatrix} w_{11}x_1 + w_{12}x_2 \\ w_{21}x_1 + w_{22}x_2 \end{pmatrix} = \begin{bmatrix} 0.22 \\ 0.26 \end{bmatrix}$$

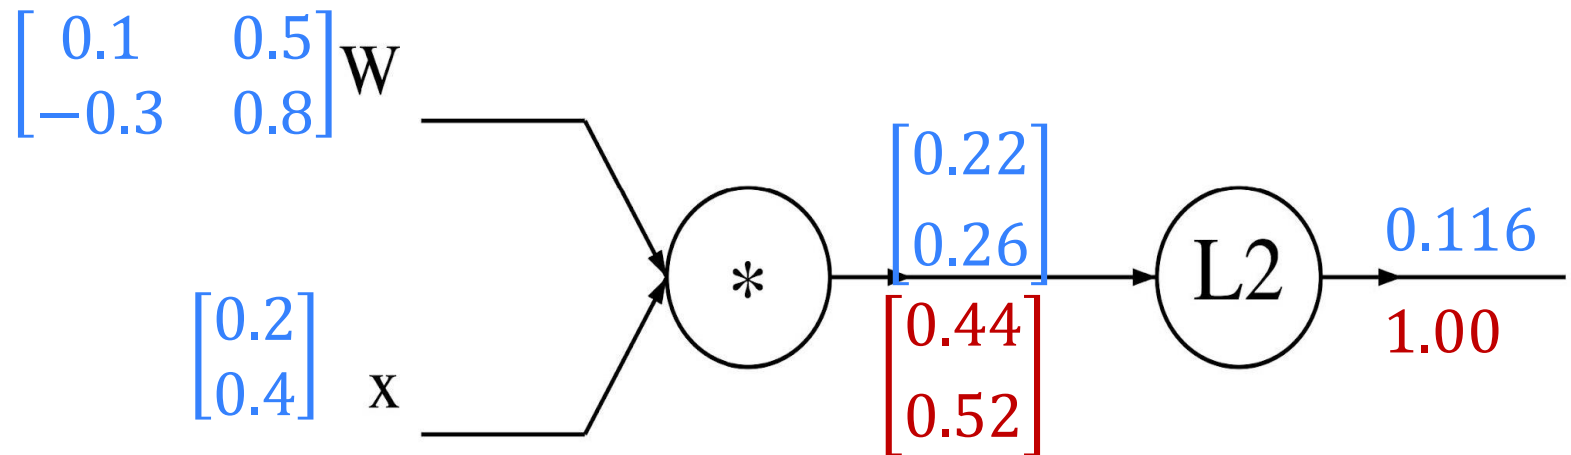
$$f(q) = \sum_{i=1}^2 q_i^2 = 0.116$$



A vectorized example: $f(\mathbf{x}, \mathbf{W}) = \sum_{i=1}^n (\mathbf{W} \cdot \mathbf{x})_i^2$

Backward:

$$f(q) = \sum_{i=1}^2 q_i^2 \longrightarrow \frac{\partial f}{\partial q_i} = 2q_i \longrightarrow \boxed{\nabla_q f = 2q}$$

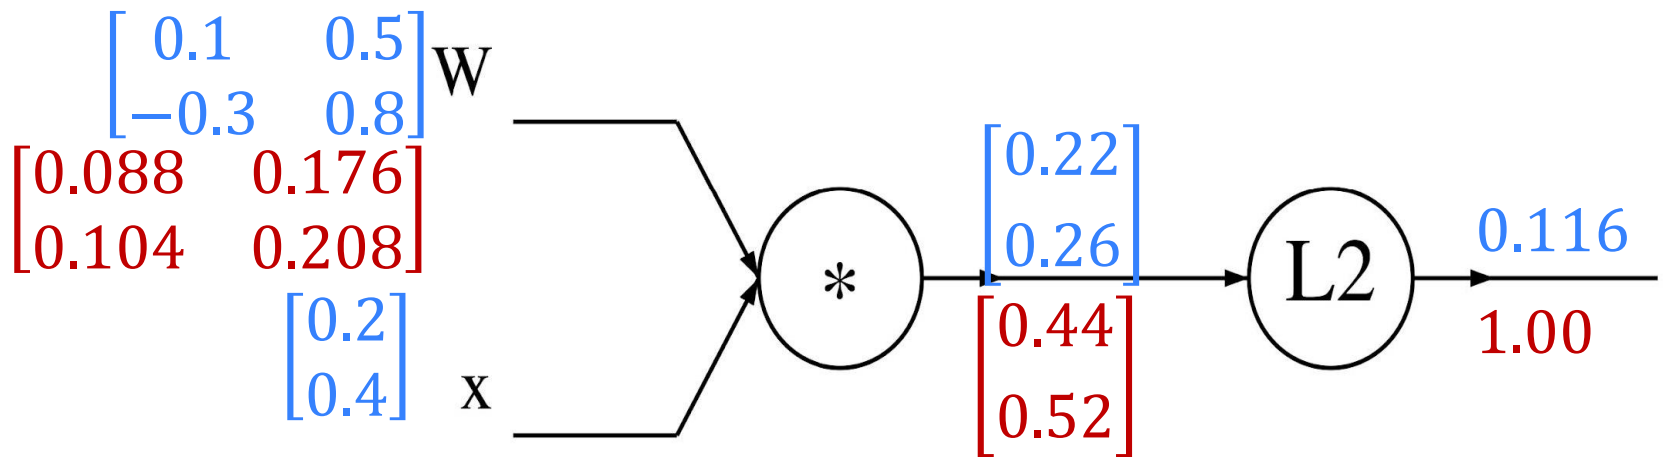


A vectorized example: $f(x, W) = \sum_{i=1}^n (W \cdot x)_i^2$

Backward:

$$q = Wx = \begin{pmatrix} w_{11}x_1 + w_{12}x_2 \\ w_{21}x_1 + w_{22}x_2 \end{pmatrix} \longrightarrow \begin{array}{ll} \frac{\partial q}{\partial w_{11}} = x_1 & \frac{\partial q}{\partial w_{12}} = x_2 \\ \frac{\partial q}{\partial w_{21}} = x_1 & \frac{\partial q}{\partial w_{22}} = x_2 \end{array}$$

$$\frac{\partial f}{\partial w_{ij}} = \sum_k \frac{\partial f}{\partial q_k} \frac{\partial q_k}{\partial w_{ij}} = 2q_i x_j \longrightarrow \boxed{\nabla_W f = 2qx^T}$$



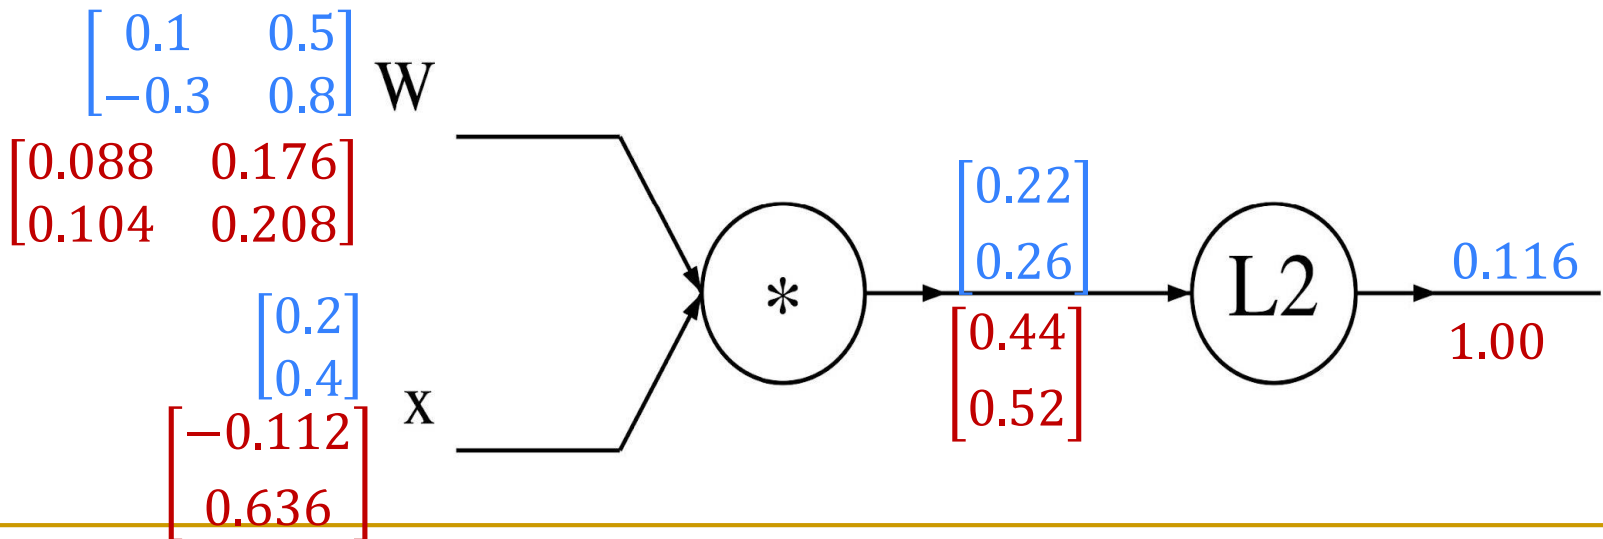
A vectorized example: $f(\mathbf{x}, \mathbf{W}) = \sum_{i=1}^n (\mathbf{W} \cdot \mathbf{x})_i^2$

Backward:

$$\mathbf{q} = \mathbf{W}\mathbf{x} = \begin{pmatrix} w_{11}x_1 + w_{12}x_2 \\ w_{21}x_1 + w_{22}x_2 \end{pmatrix}$$

$$\begin{aligned} \frac{\partial q}{\partial x_1} &= \begin{pmatrix} w_{11} \\ w_{21} \end{pmatrix} \\ \frac{\partial q}{\partial x_2} &= \begin{pmatrix} w_{12} \\ w_{22} \end{pmatrix} \end{aligned}$$

$$\frac{\partial f}{\partial x_i} = \sum_k \frac{\partial f}{\partial q_k} \frac{\partial q_k}{\partial x_i} = \sum_k 2 q_k W_{ki} \quad \longrightarrow \quad \boxed{\nabla_{\mathbf{x}} f = \mathbf{W}^T 2\mathbf{q}}$$



General tips for computation

- Derive error signal (upstream gradient) directly, avoid explicit computation of huge local derivatives
- Write out expression for a single element of the Jacobian, then deduce the overall formula
- Keep consistent indexing conventions, order of operations
- Use dimension analysis

Outline

- ❖ Modeling one neuron
 - ❖ Activation functions
 - ❖ Fully connected feed-forward network
 - ❖ How to train a multi-layer network
 - ❖ Representational power of NN
-

Representational power

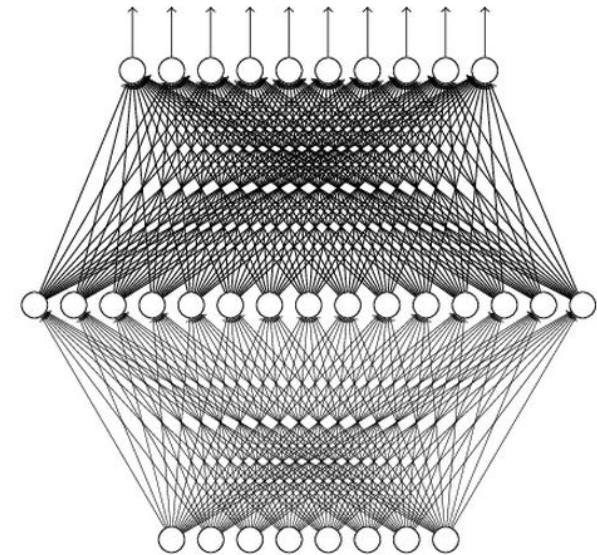
- ❖ Neural Networks with fully-connected layers define a family of functions that are parameterized by the weights of the network.
- ❖ A Neural Network with at least one hidden layer are *universal approximators*, which means that it can approximate any continuous function.

Universality Theorem

Any continuous function f

$$f : R^N \rightarrow R^M$$

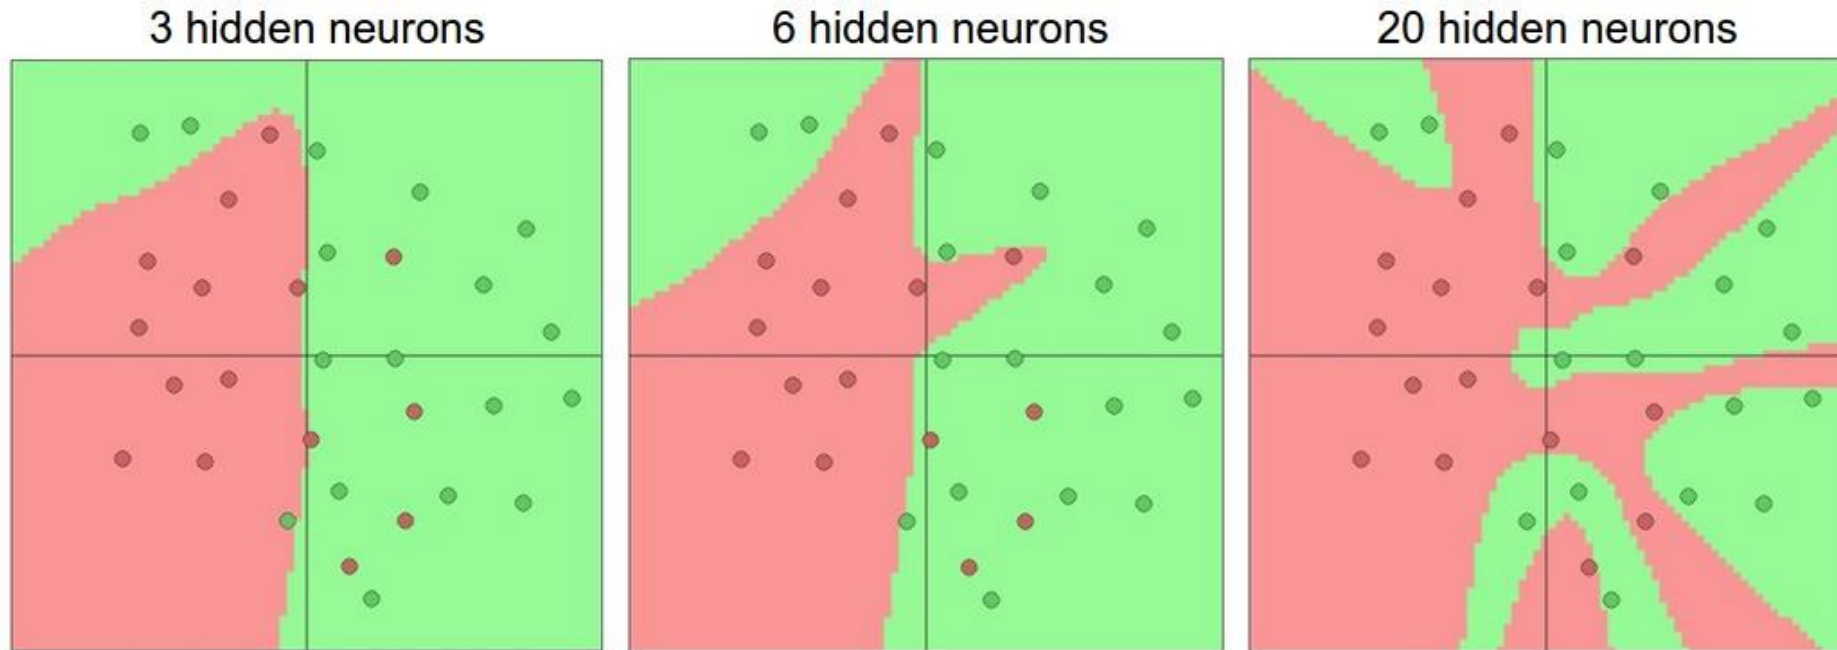
can be realized by a network
with one hidden layer
(given **enough** hidden neurons)



Reference for the reason:

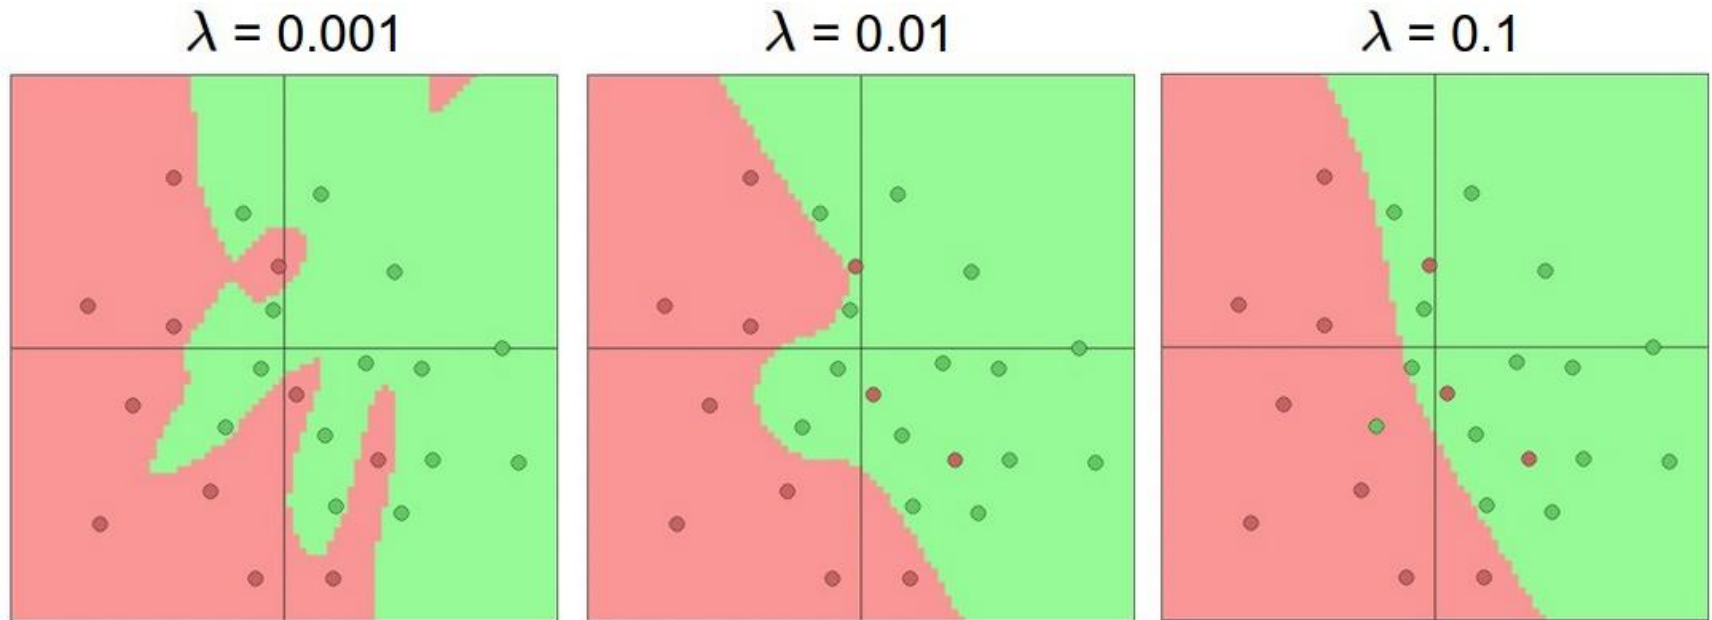
<http://neuralnetworksanddeeplearning.com/chap4.html>

Setting number of layers and their sizes

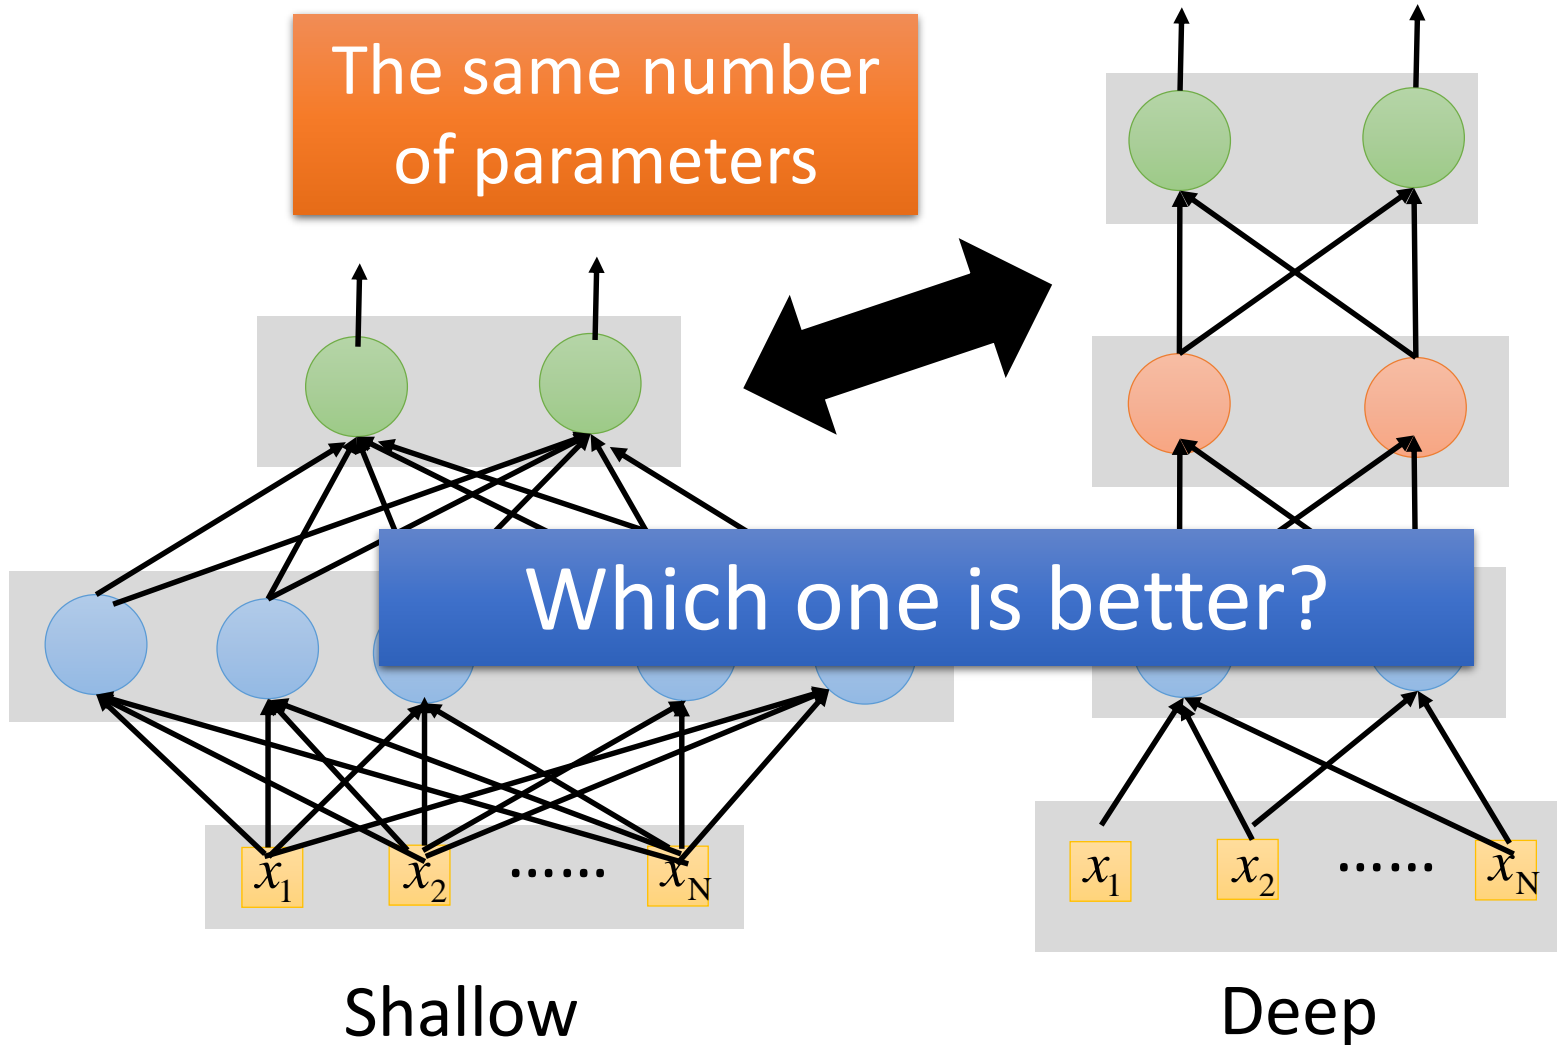


- ❖ Neural Networks with more neurons can express more complicated functions.
- ❖ But a model with high capacity fits the noise in the data instead of the (assumed) underlying relationship : **overfitting**.

- ❖ Use as big of a neural network as your computational budget allows, and use other regularization techniques to control overfitting.



Fat + Short vs. Thin + Tall



Thin + Tall vs. Fat + Short

Layer X Size	Word Error Rate (%)	Layer X Size	Word Error Rate (%)
1 X 2k	24.2		
2 X 2k	20.4		
3 X 2k	18.4		
4 X 2k	17.8		
5 X 2k	17.2	1 X 3772	22.5
7 X 2k	17.1	1 X 4634	22.6
		1 X 16k	22.1

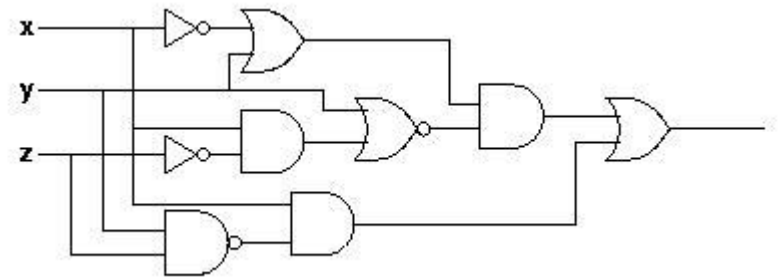
Why?

Analogy

Logic circuits

- Logic circuits consists of **gates**
- **A two layers of logic gates** can represent **any Boolean function**.
- Using multiple layers of logic gates to build some functions are much simpler

➡ less gates needed

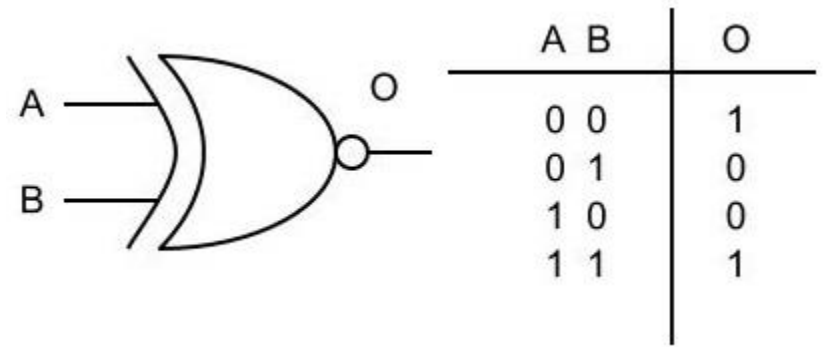


Neural network

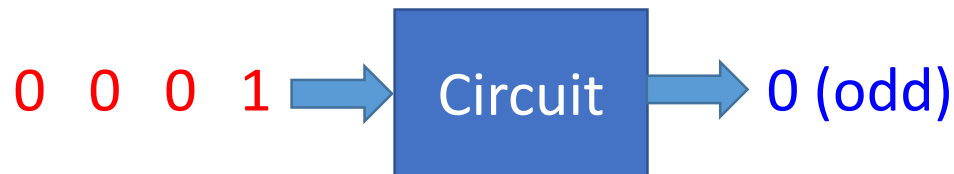
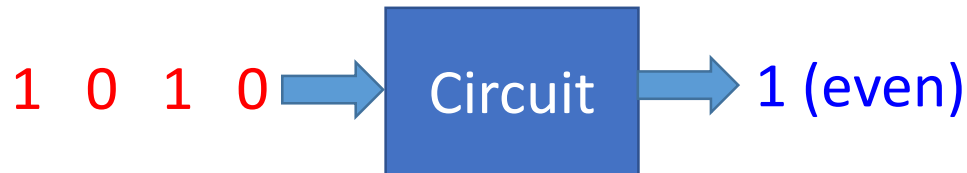
- Neural network consists of **neurons**
- **A hidden layer network** can represent **any continuous function**.
- Using multiple layers of neurons to represent some functions are much simpler

➡ less parameters ➡ less data?

Analogy

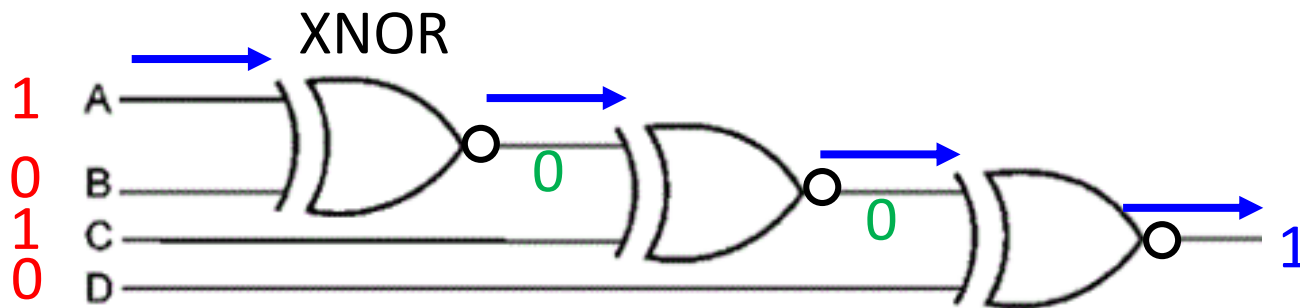


- E.g. parity check



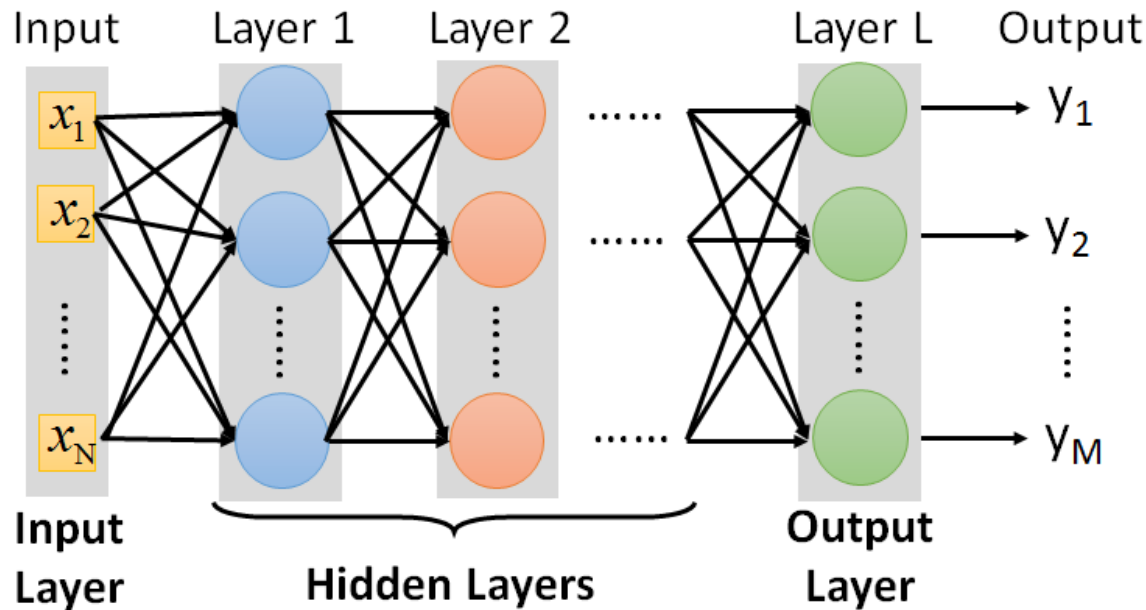
For input sequence with d bits,

Two-layer circuit need $O(2^d)$ gates.



With multiple layers, we need only $O(d)$ gates.

Design the Network



- How many layers? How many neurons for each layer?

Trial and Error

+

Intuition

课程部分材料来自他人
和网络，仅限教学使用
，请勿传播，谢谢！